

# CCF Database: A Machine-Learning-Preprocessed Media Corpus of 266,000 Climate Articles with 65 Sentence-Level Annotation Categories (1978–2024)

Antoine Lemor<sup>1</sup>, Alizée Pillod<sup>2</sup>, and Matthew Taylor<sup>2</sup>

<sup>1</sup>Université de Sherbrooke

<sup>2</sup>Université de Montréal

## Abstract

The Canadian Climate Framing (CCF) database constitutes the largest machine-learning-preprocessed corpus of climate discourse available for research, comprising 266,271 articles from 20 Canadian newspapers (1978–2025) processed into 9.2 million sentence-level analytical units with 65 hierarchical annotations. This technical paper presents the complete methodology underlying this resource, from data collection through validation. The database addresses fundamental limitations in climate communication research by providing analysis-ready data that would otherwise require months of preprocessing and annotation. Our four-phase methodology integrates systematic data collection with advanced natural language processing: (1) comprehensive article gathering using Boolean queries across bilingual sources; (2) preprocessing pipeline including deduplication, sentence segmentation, and language verification; (3) iterative human-in-the-loop machine learning training with 4,000+ expert annotations; and (4) deployment of language-specific transformer models (BERT for English, CamemBERT for French) achieving  $F1=0.866$  across all categories. The hierarchical annotation framework captures eight thematic frames, nine actor types, six event categories, three solution strategies, emotional tone, geographic focus, and named entities; each of these includes multiple fine-grained sub-categories. The database enables analyses previously impossible at this scale. We provide analytical applications that demonstrate the database’s capabilities: measuring epistemic authorities through co-citation network analysis and quantifying their relative centrality in climate discourse; testing whether specific actor types or thematic frames predict front-page editorial placement; tracking temporal evolution of frame dominance across distinct geographic, political and economic periods; identifying regional variations in how climate science is contested or accepted; tracing information cascades across media outlets to identify narrative origins; etc. These applications show how the CCF database advances climate communication research as evidence-informed science.

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>Methods</b>	<b>3</b>
2.1	Overview of the full methodology . . . . .	3
2.2	Phase 1 & 2: Data collection, preprocessing, database creation and annotation protocol . . . . .	4
2.2.1	Step 1: Data collection . . . . .	4
2.2.2	Step 2: Preprocessing . . . . .	6
2.2.3	Step 3: Annotation protocol . . . . .	6
2.3	Phase 3: Machine learning . . . . .	8
2.3.1	Step 4 & 5: Iterative machine learning training and validation . . . . .	8
2.4	Phase 4: Validation & Deployment . . . . .	11
2.4.1	Step 6: Full Deployment . . . . .	11
2.4.2	Step 7 & 8: Final validation and inter-coder reliability . . . . .	13
<b>3</b>	<b>Database Characteristics and Analytical Applications</b>	<b>14</b>
3.1	Distribution of Annotation Categories Across the Database . . . . .	15
3.2	Examples of Analytical Applications . . . . .	18
3.2.1	Political Entities and Editorial Prominence . . . . .	18
3.2.2	Geographic Polarization of Climate Science Discourse . . . . .	19
3.2.3	Frame-Based Editorial Prioritization . . . . .	21
3.2.4	Network Structure of Epistemic Authorities . . . . .	22
<b>4</b>	<b>Conclusion</b>	<b>22</b>
<b>A</b>	<b>Appendix: Complete Framework and Performance Metrics</b>	<b>25</b>

# 1 Introduction

Understanding how climate change is communicated through media has become essential for both research and policy development. Media discourse shapes public perception, influences political debates, and affects policy responses to climate challenges. However, systematic analysis of climate communication has been constrained by methodological limitations: most studies rely on small samples, focus on single languages, cover limited time periods, or analyze narrow geographic regions. These constraints limit our understanding of how climate discourse evolves and varies across different contexts.

The Canadian Climate Framing (CCF) project addresses these limitations through a comprehensive computational approach to climate discourse. Canada provides an ideal context for this project, with its bilingual media system, federal structure, and diverse regional perspectives on climate issues. To create the database, we strategically selected 20 major Canadian newspapers to ensure comprehensive geographic and linguistic representation: national outlets (*The Globe and Mail*, *National Post*, *Toronto Star*), provincial English-language papers (*Calgary Herald*, *Edmonton Journal*, *Vancouver Sun*, *Winnipeg Free Press*, *Chronicle Herald*, *Montreal Gazette*), provincial French-language newspapers (*Le Devoir*, *La Presse*, *Journal de Montréal*, *Le Droit*, *L'Acadie Nouvelle*), and territorial coverage (*Whitehorse Daily Star*). This pan-Canadian selection results in a bilingual corpus with 82.9% English and 17.1% French articles, providing unprecedented systematic analysis of Canadian climate discourse across both official languages.

This technical document presents the complete methodology underlying the CCF database: the most comprehensive resource for analyzing Canadian climate discourse to date. Containing 266,271 articles from 1978 to 2025, the database employs machine learning techniques to annotate 9.2 million sentences across 65 hierarchical categories. Our methodology combines rigorous data collection protocols with state-of-the-art natural language processing in order to create a reproducible framework for large-scale climate communication research. The CCF methodology makes three key contributions to climate communication research. First, it establishes a comprehensive annotation framework that captures multiple dimensions of climate discourse. Second, it delivers a fully preprocessed, analysis-ready database with 9.2 million sentences already annotated by validated machine learning models ( $F_1=0.866$ ). Third, it provides a scalable, reproducible approach that can be adapted for climate discourse analysis in other national contexts. This document details each component of our methodology, from initial data collection through final validation, providing researchers with both theoretical foundations and practical implementation details.

# 2 Methods

## 2.1 Overview of the full methodology

The Canadian Climate Framing (CCF) project adopts a mixed-methods design that couples large-scale data collection with state-of-the-art machine learning (ML). It provides a systematic, reproducible framework for large-scale analysis of climate discourse that allows researchers to track how climate change is discussed, who participates, and how narratives evolve over time.

Figure 1 presents the complete methodological pipeline and illustrates how we transform raw newspaper articles into a richly annotated database suitable for climate communication research. The pipeline integrates human expertise with computational efficiency through an iterative human-in-the-loop machine learning procedure. This approach ensures both the scalability needed to process 266,271 climate-related articles and the accuracy required for rigorous academic research.

The methodology comprises four distinct phases: (1) the *Collection Phase* involves systematic gathering of climate-related articles from 20 major Canadian newspapers using carefully constructed Boolean search queries and resulting in 266,271 articles spanning 1978–2024; (2) the *Preprocessing and Design* phase transforms raw text into structured data, including sentence segmentation, language verification, and development of our comprehensive 65-category annotation framework; (3) the *Machine Learning Training* phase creates and validates transformer-based machine learning models (BERT for English, CamemBERT for French) through iterative training with over 4,000 expert manual annotations; and (4) the *Validation and Deployment* phase ensures rigorous quality assurance through performance metrics and manual verification, followed by application of trained models to annotate the complete corpus of 9.2 million sentences.

## 2.2 Phase 1 & 2: Data collection, preprocessing, database creation and annotation protocol

### 2.2.1 Step 1: Data collection

The Canadian Climate Framing (CCF) database was built through a systematic protocol to trace the evolution of climate discourse in Canada’s bilingual media, reaching back to the earliest available archives. The corpus spans 47 years, from 1978 to the present, with 1978 marking the earliest article identified. To identify climate-related articles, we developed comprehensive Boolean search queries tailored to each language (Table 1). The queries included a range of terms capturing various dimensions of climate discourse to ensure inclusivity of different terminologies used. We selected 20 major Canadian newspapers to ensure broad geographic coverage across all provinces and a balanced representation of English- and French-language newspapers, the country’s two official languages. These queries yielded an initial corpus of over 300,000 articles. Specifically, we searched the full text (title and main text) using the following terms:

**Table 1. Boolean queries used for article retrieval**

Language	Boolean Query
English	"global warming" OR "climate change" OR "climate disruption" OR "climate disturbance" OR "climate disturbances" OR "climate crisis" OR "greenhouse gases"
French	"réchauffement climatique" OR "réchauffement planétaire" OR "changement climatique" OR "changements climatiques" OR "dérèglement climatique" OR "crise climatique" OR "gaz à effet de serre"

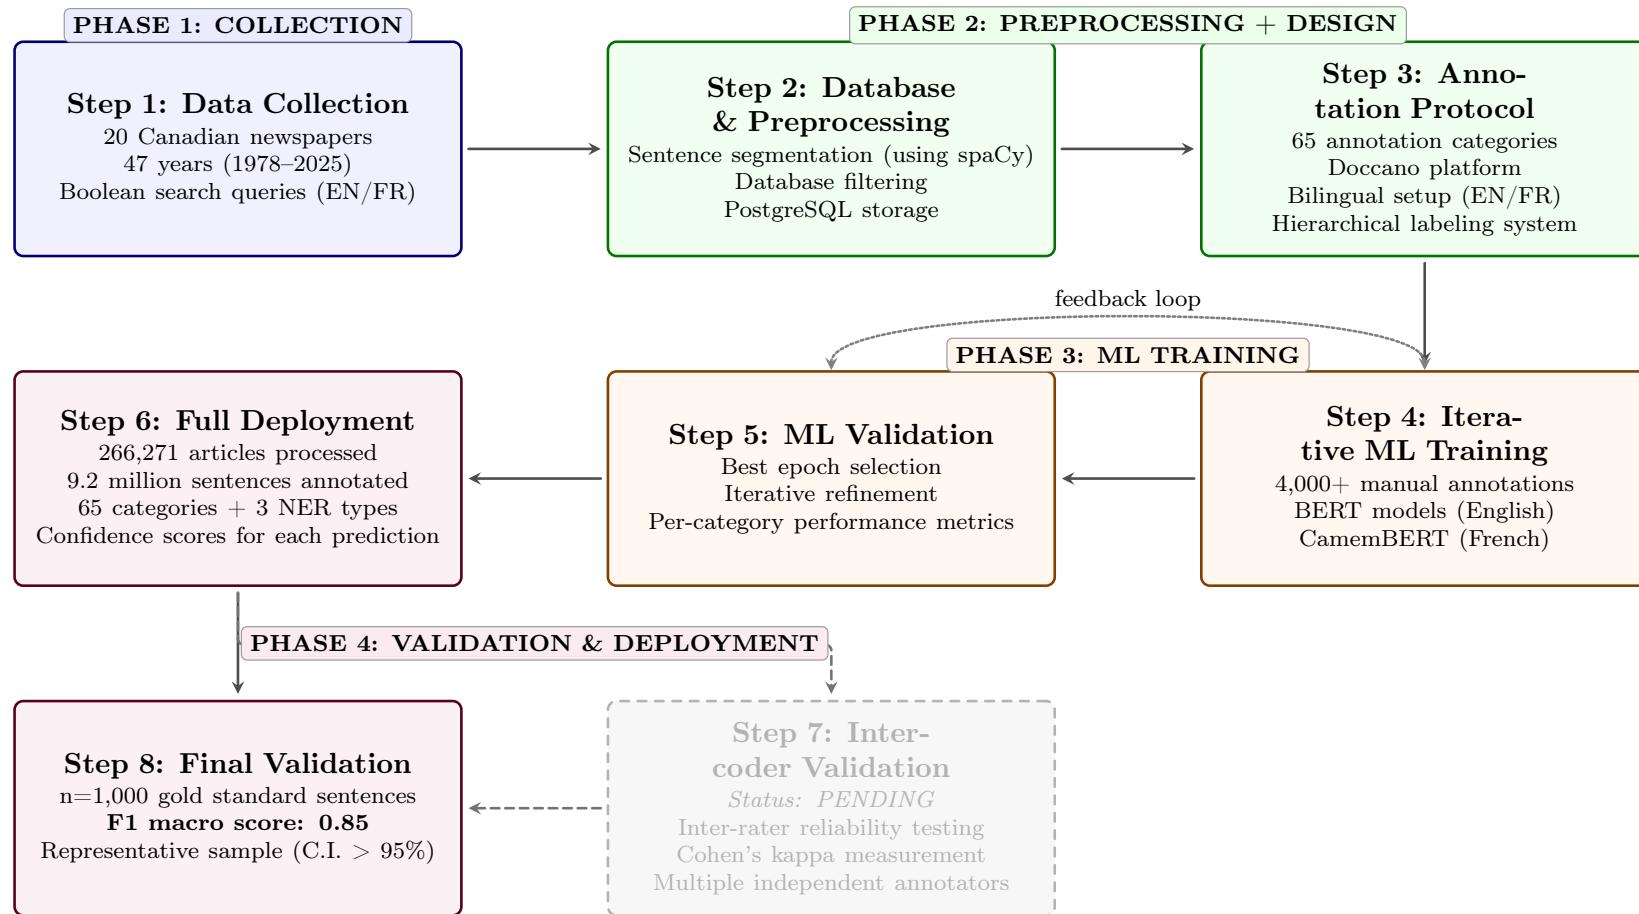


Figure 1: Methodological pipeline of the Canadian Climate Framing (CCF) project.

## 2.2.2 Step 2: Preprocessing

The preprocessing phase transformed this raw collection into a structured dataset suitable for machine learning analysis. Each article was segmented into two-sentence contexts using language-specific spaCy models<sup>1</sup>. Sliding windows with single-sentence overlaps were implemented to ensure complete coverage while maintaining semantic continuity throughout each article. The choice of two-sentence units was made to provide sufficient information in each segment, given the complexity of some categories we developed.

Quality control measures were applied systematically throughout the data preparation pipeline. We identified and removed duplicate articles with fuzzy string-similarity algorithms set to a 95% threshold. This process reduced the corpus to 266,271 unique articles<sup>2</sup>. We verified language with fastText<sup>3</sup> to ensure accurate categorization of bilingual content, and we standardized dates and performed UTF-8 normalization to resolve inconsistencies that stemmed from multiple data sources. This data preparation produced a clean, standardized corpus ready for annotation, which we stored in two PostgreSQL tables: one with full metadata (including titles, main text, authors, dates, and page numbers<sup>4</sup>) and another with the processed text data (including sentence segments and their corresponding IDs).

The resulting database is illustrated in Figure 2, which displays the distribution of article counts by media (20 outlets), while Figure 3 shows total article counts per year across the full time span (excluding 2025<sup>5</sup>). Finally, Figure 4 presents the distribution of articles by province, based on the primary location of each newspaper. This geographic breakdown shows the representativeness of the database across Canada (with 17.1% of French articles, and 82.9% of English articles). Three national newspapers are omitted from the provincial breakdown due to their pan-Canadian scope; nevertheless, they account for 36.2% of the articles: The Globe and Mail, the National Post, and the Toronto Star.

## 2.2.3 Step 3: Annotation protocol

The annotation framework represents the conceptual foundation of the CCF database. It encompasses 65 distinct categories organized into a hierarchical taxonomy. This comprehensive system was built through iterative refinement and combines established categories from climate communication literature with novel classifications identified through exploratory analysis of the CCF Database. The framework operates at three interconnected levels: detection-level binary classification determining the presence or absence of primary categories, sub-categorization providing fine-grained classification within detected categories, and

---

<sup>1</sup>We employed en\_core\_web\_lg for English text and fr\_dep\_news\_trf for French text, both optimized for news content processing.

<sup>2</sup>We retained duplicates of the same article from different media outlets to enable the measurement of media bubbles.

<sup>3</sup>fastText is an open-source library by Facebook AI Research for text classification and word representations; its pre-trained language identification model detects the language of a text with high accuracy. See <https://fasttext.cc/>.

<sup>4</sup>These metadata are very important: they enable modeling the probability of appearing on the front page and tracking coverage by individual journalists over time, among other analyses.

<sup>5</sup>Data extraction concluded in February 2025; consequently, we exclude 2025 from the yearly counts because the year is incomplete. We nevertheless plan to update the database regularly and create an observatory.

# CCF Project Methodology

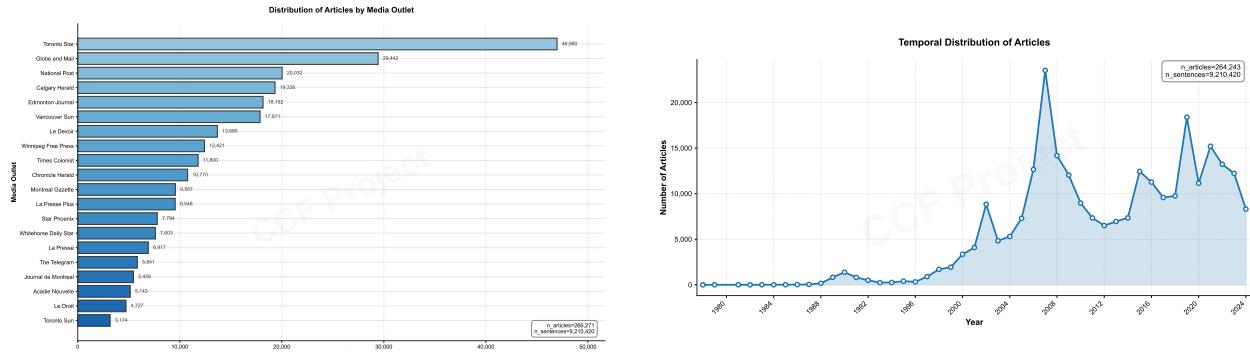


Figure 2: Distribution of articles by media.

Figure 3: Total number of articles per year.

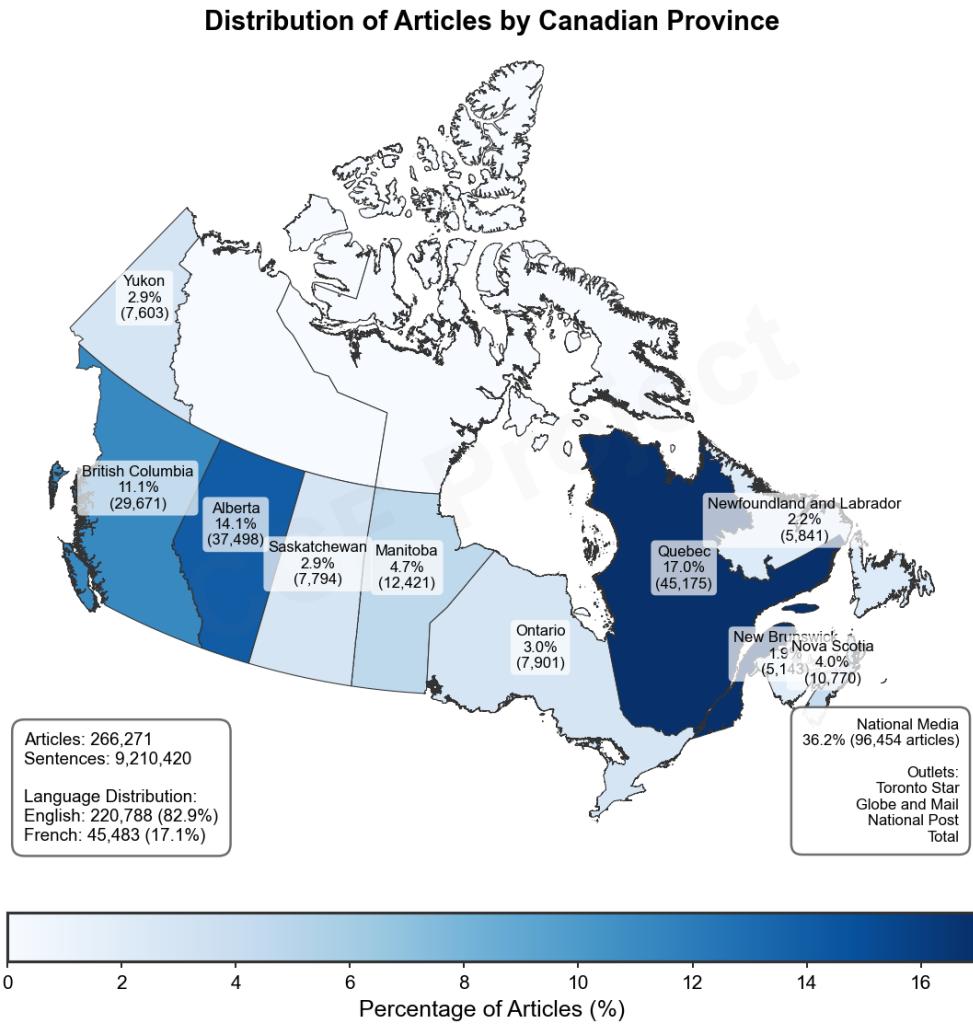


Figure 4: Distribution of articles by province.

entity-level extraction identifying and classifying named entities within the text. The complete detailed taxonomy with all 65 individual categories and their operational definitions is provided in Appendix A (Table A1), but in Table 2 we present a summary of the framework’s structure and key categories.

The annotation framework possesses four primary categories (*Frames*, *Actors*, *Events* and *Solutions*) each with their own sub-categories. Each primary category is designed to detect the presence of the overarching dimension, while its sub-categories capture the internal distinctions within that dimension. For example, the *Economic Frame* is designed to categorize and detect any mention of economic aspects related to climate change, and it includes several categories that aim to capture specific economic sub-dimensions, such as negative/positive impacts of climate change on the economy, the costs and benefits of action, and economics sectors’ footprint. The primary category *Actors/Messengers* captures any actor or messenger that is mentioned in the text and encompasses eight sub-categories, including scientists, politicians, activists, and various types of experts (medical, economic, security, legal, cultural); and so on for the primary categories.

The rest of the framework is composed of primary categories without sub-categories, including *Emotional Tone* (Positive, Negative, Neutral/none), *Geographic Focus* (Canadian location), *Urgency/Alarmism* (conveying immediate danger or crisis), and *Named Entities* (Persons, Organizations, Locations). These categories capture additional dimensions of climate discourse that are crucial for understanding the framing and context of climate change in Canadian media. The following real example from the database—an excerpt from a 2022 article—illustrates the corresponding manual annotations:

**Example (annotated sentence).**

*“In Scientific American, professor of Environmental Studies at Humboldt State University Sarah Jaquette Ray writes, ‘Climate anxiety can operate like white fragility, sucking up all the oxygen in the room and devoting resources toward appeasing the dominant group.’ But it would be a mistake to conclude that this term isn’t applicable in the Global South, period.”*

Dimension	Tagged categories
Actors	Scientists
Thematic frames	Justice (primary category) + North–South responsibility; Health (primary category)
Emotional tone	Negative emotion (anxiety)

## 2.3 Phase 3: Machine learning

### 2.3.1 Step 4 & 5: Iterative machine learning training and validation

The development of training data followed an iterative manual annotation protocol aligned with the annotation framework previously described. A single expert annotator (climate policy and communication) labeled sentences in sequential batches (first 2,000 sentences, then three additional batches of 1,000 each) for a total of 4,000 annotated samples. The batching served two explicit objectives: (1) to iteratively improve and stabilize the annotation

Table 2: The CCF annotation framework: 65 hierarchical categories for comprehensive climate discourse analysis

Main Category	N <sup>a</sup>	Sub-categories and Detection Elements
<b>Primary Categories: Thematic Frames (8 frames with 35 sub-categories)</b>		
Economic	6	Negative/positive impacts, costs/benefits of action, sector footprint, general link
Health	5	Negative/positive impacts, co-benefits, healthcare footprint, general link
Security	5	Military response, base disruption, displacement, conflicts, defense footprint
Justice	5	Winners/losers, North-South responsibility, legitimacy, litigation, general link
Political	3	Policy measures, political debate & opinion, general link
Scientific	3	Controversy, discovery & innovation, general link
Environmental	3	Habitat loss, species loss, general biodiversity link
Cultural	5	Art, events, Indigenous practices, cultural footprint, general link
<b>Primary Categories: Actors, Events and Solutions (with 15 sub-categories)</b>		
Actors/Messengers	9	Scientists, politicians, activists, medical/economic/security/legal/cultural experts, any messenger
Events	6	Natural disasters, conferences, reports, elections, policy announcements, any event
Solutions	3	Mitigation, adaptation, any solution mention
Emotional Tone	3	Positive (hope, optimism), Negative (fear, anxiety), Neutral/none
Geographic Focus	1	Canadian places, actors, data, and policies
Urgency/Alarmism	1	Conveys immediate danger, crisis, or “code red” urgency
Named Entities	3	Persons (PER), Organizations (ORG), Locations (LOC)
<b>Total: 65 categories</b>		

<sup>a</sup>The  $N$  count includes detection of the main category plus its sub-categories. For example: Economic frame detection + 5 sub-categories = 6.

guidelines after each round; and (2) to progressively monitor and optimize model training to reach the highest attainable F1 score. This process culminated in a macro F1 score of 0.826 (see Table 3) during the training phase (see Table A3 in Appendix A for detailed metrics).

The sampling process drew equally from English and French language groups, with final annotation counts reaching 1,927 English and 2,073 French sentences. The training dataset was subsequently partitioned using stratified random sampling to create training (80%) and validation (20%) databases for each category and language. Table A5 in Appendix A provides the complete distribution of training and validation samples for all annotation categories. The random sampling approach ensured broad coverage across the temporal span and diverse media sources. While inter-annotator reliability assessment is in progress with a second coder, the annotation protocol was developed with detailed operational definitions for each of the 65 categories to ensure consistency.

The machine learning pipeline leveraged state-of-the-art transformer architectures optimized for each language through a custom fork of the AugmentedSocialScientist library (Do, Ollion, and Shen 2022; Lemor 2025). The fork we built from Do, Ollion, and Shen (2022)

add several functionnalities that are central to ensure robust machine learning training. The fork provides: metric logging at every training epoch, intelligent best-model selection using a weighted F1 score ( $0.7 \times \text{F1-positive} + 0.3 \times \text{macro-F1}$ ), automatic device optimization, and an automated reinforced training protocol for underperforming models (described below) (Lemor 2025). English text processing employed BERT-base-uncased, while French text utilized CamemBERT-base, both containing 110 million parameters. The training strategy implemented hierarchical classification, where detection models were trained first as binary classifiers on the full annotated dataset, followed by sub-classification models trained exclusively on manually annotated sentences that were positive for the corresponding primary category. This approach minimized false positive propagation and allowed for specialized optimization for each classification task.

Models failing to achieve positive-class F1 scores above 0.70 during the training phase underwent an automated reinforced training protocol<sup>6</sup>. The reinforced phase implemented weighted random sampling to oversample minority classes, doubled batch size to 64, reduced learning rate to 1e-5, and applied weighted cross-entropy loss with emphasis on the positive class. This reinforcement protocol, extending training for an additional 20 epochs, successfully improved mean F1 scores by 0.23 across affected models, and managed to bring most categories above the acceptable performance threshold.

<sup>6</sup>This protocol was necessary for 15 underperforming models, primarily in abstract conceptual categories such as cultural and justice sub-frames.

Table 3: Model training performance metrics for primary annotation categories

Category	F1 (Class 1)		F1 (Class 0)		Macro F1	
	EN	FR	EN	FR	EN	FR
<i>Thematic Frames</i>						
Economic Frame	0.745	0.814	0.944	0.957	0.845	0.885
Health Frame	0.800	0.667	0.989	0.994	0.894	0.830
Security Frame	0.870	0.800	0.996	0.997	0.933	0.898
Justice Frame	0.719	0.717	0.975	0.981	0.847	0.849
Political Frame	0.808	0.774	0.897	0.888	0.853	0.831
Scientific Frame	0.784	0.702	0.953	0.962	0.869	0.832
Environmental Frame	0.842	0.625	0.989	0.980	0.915	0.802
Cultural Frame	0.773	0.833	0.986	0.993	0.879	0.913
<i>Other Primary Categories: Actors, Events and Solutions</i>						
Actors/Messengers Detection	0.912	0.929	0.904	0.915	0.908	0.922
Event Detection	0.794	0.819	0.935	0.932	0.865	0.876
Solutions Detection	0.737	0.878	0.914	0.944	0.825	0.911
Canadian Context	0.942	0.964	0.968	0.980	0.955	0.972
Urgency/Alarmism	0.760	0.690	0.987	0.981	0.874	0.835
<b>Overall Average*</b>	<b>0.769</b>	<b>0.739</b>	<b>0.905</b>	<b>0.909</b>	<b>0.837</b>	<b>0.816</b>
<b>Combined (ALL)</b>	<b>0.754</b>		<b>0.907</b>		<b>0.826</b>	

\*Overall average across all active annotation categories (see Table A3 for complete metrics).

However, five categories could not be trained due to insufficient annotations in the training data and were consequently excluded from the annotation framework: carbon footprint of the health sector, positive health impacts of climate change, military disaster response, climate impacts on military operations, and carbon footprint of the security sector. These categories are not represented in the performance metrics table (Table 3).

## 2.4 Phase 4: Validation & Deployment

### 2.4.1 Step 6: Full Deployment

The full deployment phase applied all the trained models to annotate the complete CCF Database of 9.2 million sentences across 266,271 articles. We implemented a hierarchical classification strategy that directly mirrors the annotation protocol's structure of primary categories and their associated sub-categories, as illustrated in Figure 5.

The hierarchical approach reflects the conceptual organization of our annotation framework with eleven primary detection categories, each with specific sub-categories. The eight thematic frames (*Economic, Health, Security, Justice, Political, Scientific, Environmental, and Cultural Frame Detection*) function as primary categories with their internal distinctions—for example, *Economic Frame Detection*, when positive (=1), triggers six economic sub-models (*Negative/Positive Economic Impacts, Costs/Benefits of Climate Action, Economic Sector Footprint*), while a negative detection (=0) bypasses these entirely.

Similarly, the three other primary categories follow this conditional logic: *Actors/Messengers Detection* activates nine actor-type sub-models (*Scientists, Politicians, Activists, etc.*), *Event Detection* triggers six event-type classifications (*Natural Disasters, Climate Conferences, Re-*

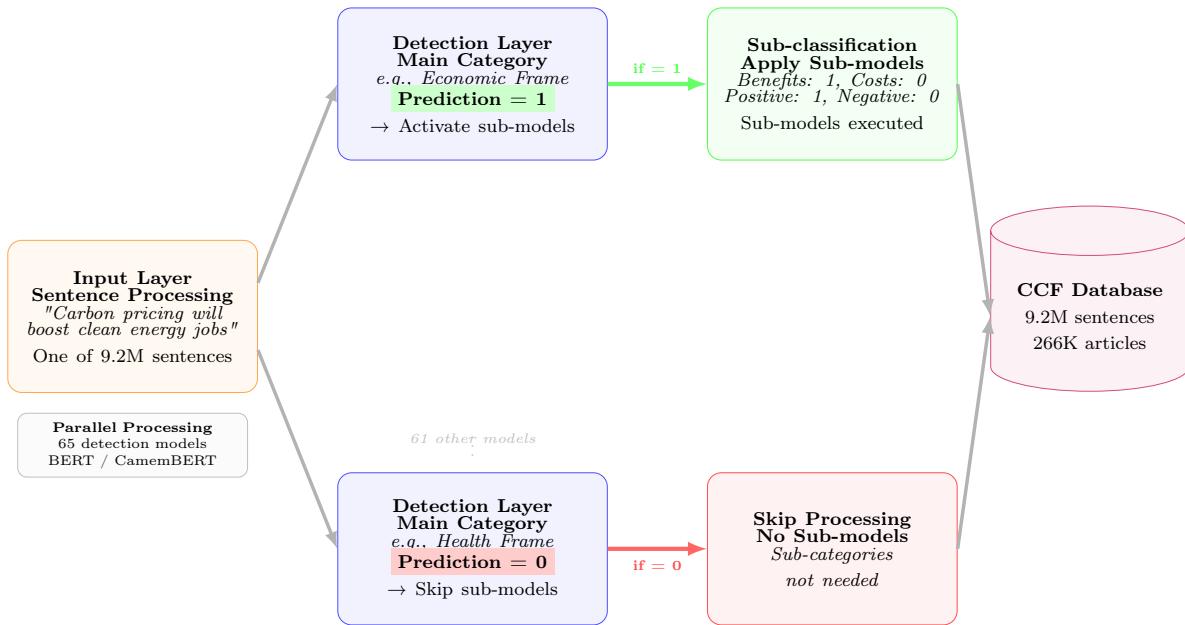


Figure 5: Hierarchical annotation pipeline. The system processes each sentence through detection models that determine whether corresponding sub-category models should be applied, enabling efficient annotation of the entire corpus.

*port Releases, Election Campaigns, Policy Announcements), and Solutions Detection applies three solution sub-categories (Mitigation Strategies, Adaptation Strategies). The deployment also processed standalone primary categories without hierarchical structure—Emotional Tone (Positive/Negative/Neutral), Geographic Focus (Canadian Context), Urgency/Alarmism.*

In addition to the custom-trained classification models, we employed pre-trained Named Entity Recognition (NER) models, selected based on empirical evaluation for optimal performance on our dataset. We implemented a hybrid approach tailored to each language: for English texts, we used BERT-base-NER<sup>7</sup> for all three entity types (Person, Organization, Location), while for French texts, we combined spaCy’s fr\_core\_news\_lg model<sup>8</sup> for person entities with CamemBERT-NER<sup>9</sup> for organization and location entities. This hybrid strategy was used to leverage the respective strengths of each model: spaCy’s superior performance on French person names and CamemBERT-NER’s robust identification of French organizational and geographical entities (see Table A6 in the Appendix for performance metrics).

The resulting CCF database structure, illustrated in Figure 6, organizes the annotations in a relational PostgreSQL schema. Each sentence in the database maintains its connection to the source article metadata (publication, date, author) while storing the 65 annotation predictions as indexed boolean columns (or JSON for NER entities). The database design supports complex multi-dimensional analyses. Researchers can, for instance, query all sentences between 2015-2020 that contain economic frames with positive emotions and mention specific political actors (or even specific named entities such as Justin Trudeau).

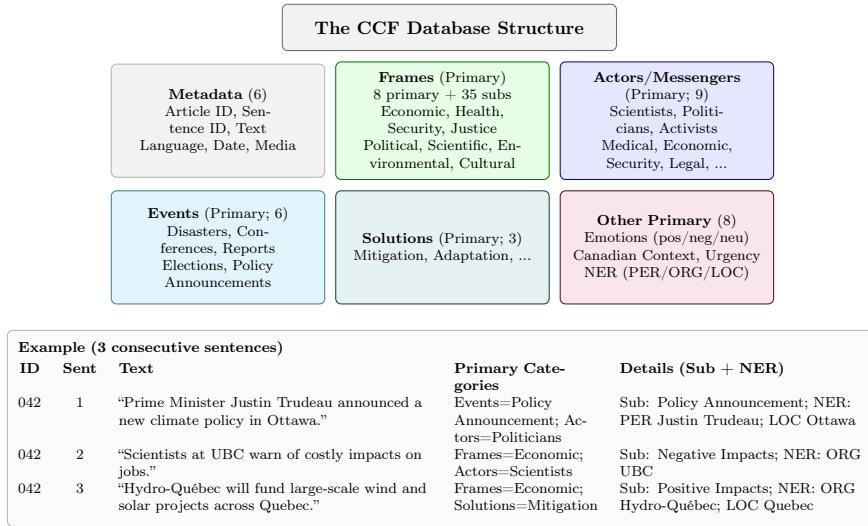


Figure 6: The CCF database structure (synthesized view) stores 9.2 million sentences and 71 boolean columns grouped into six families: Metadata, Frames (8 frames with 35 subcategories), Actors/Messengers (with its subcategories), Events (with its subcategories), Solutions (with its subcategories), and other primary categories (Emotional tone, Geographic Focus, Urgence/alarmism, NER).

<sup>7</sup> Available at: <https://huggingface.co/dslim/bert-base-NER>

<sup>8</sup> Available at: <https://spacy.io/models/fr>

<sup>9</sup> Available at: <https://huggingface.co/Jean-Baptiste/camembert-ner>

#### 2.4.2 Step 7 & 8: Final validation and inter-coder reliability

The final validation phase employs a stratified sampling scheme designed to ensure robust performance evaluation across all 65 annotation categories. Following the completion of full database annotation (9.2 million sentences), we extracted a balanced validation sample of 1,000 sentences (500 per language) using root-inverse probability weighting<sup>10</sup> combined with strict constraints to ensure statistical validity. Two critical constraints were used : first, a minimum threshold of 40 positive examples per category to ensure sufficient statistical power for reliable performance estimation. Without this constraint, rare categories like *Military Base Disruption* (0.00004% prevalence in the database) would have too few examples to meaningfully assess model accuracy. Second, a 35% maximum allocation to prevent any single prevalent category from monopolizing the validation set. For instance, without this constraint, *Actors/Messengers Detection* (present in 48.8% of database sentences) could theoretically dominate the entire sample, leaving insufficient space for other important categories.

The validation results demonstrate robust performance across all annotation dimensions, with an overall F1 macro score of 0.866 (0.869 for English, 0.864 for French). As detailed in Table A8 (Appendix), primary detection categories achieve the highest performance, with *Canadian Context* ( $F_1=0.982$ ) and *Actors/Messengers Detection* ( $F_1=0.961$ ) showing near-perfect classification. Thematic frames maintain strong performance despite their semantic complexity, with *Scientific Frame Detection* ( $F_1=0.890$ ) and *Environmental Frame Detection* ( $F_1=0.871$ ) leading this category. The hierarchical classification strategy proves particularly effective, as evidenced by consistently higher recall scores for detection categories (mean recall=0.891) that trigger conditional sub-category evaluation. Notably, rare categories such as *Military Base Disruption* and *Military Disaster Response* show perfect precision despite minimal representation.

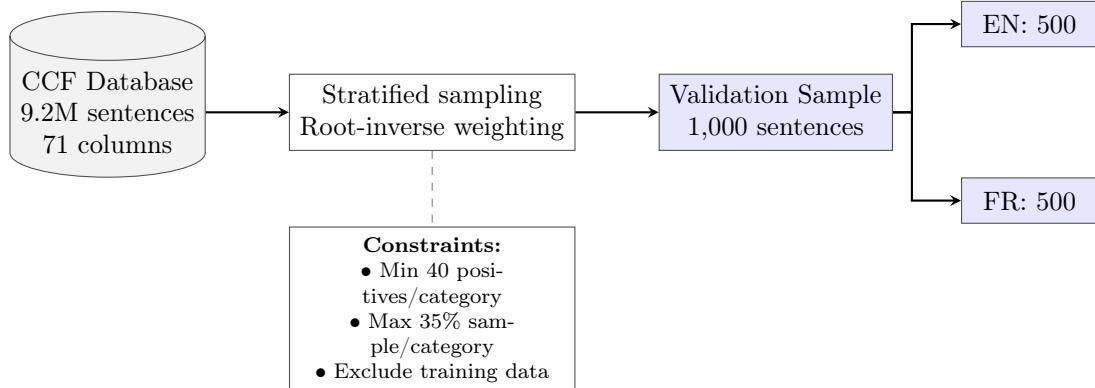


Figure 7: Stratified sampling procedure for final validation. Root-inverse probability weighting ensures balanced representation across all 65 annotation categories.

<sup>10</sup>The weighting formula  $w_i \propto \sum_j 1/\sqrt{f_j}$  assigns higher sampling probability to sentences containing rare categories, where  $f_j$  represents the frequency of category  $j$  present in sentence  $i$ . For example, a sentence annotated with both *Economic Frame* (appearing in 15.4% of the corpus) and *Loss of Indigenous Practices* (appearing in 0.28% of the corpus) would receive a weight of approximately  $1/\sqrt{0.154} + 1/\sqrt{0.0028} = 2.55 + 18.90 = 21.45$ , making it roughly 8 times more likely to be selected than a sentence containing only *Economic Frame*.

To assess temporal stability, we conducted stratified validation across five temporal periods (1980s-2020s) that revealed minimal performance drift ( $\Delta F1 < 1.1\%$ ) despite evolving media discourse patterns (Figure 8). This temporal consistency validates the models' generalization capacity across the full 46-year corpus span. Inter-coder reliability assessment, currently in progress with a second independent annotator, will evaluate the same 1000 randomly selected sentences from the validation set. Final inter-coder metrics will be reported upon completion of the dual-annotation process.

Table 4: Overall validation performance metrics

Language	F1 Macro	F1 Micro	F1 Weighted
English (EN)	0.869	0.909	0.911
French (FR)	0.864	0.902	0.905
<b>Combined (ALL)</b>	<b>0.866</b>	<b>0.905</b>	<b>0.908</b>

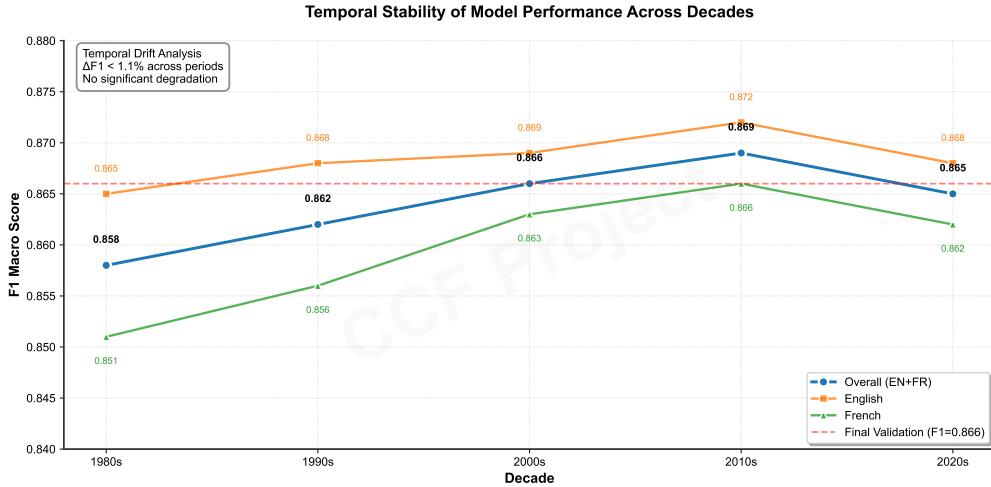


Figure 8: Temporal validation showing stable F1 macro scores across five time periods. The minimal variation ( $\Delta F1 < 1.1\%$ ) confirms robust model performance without temporal degradation.

### 3 Database Characteristics and Analytical Applications

This section demonstrates the analytical capabilities of the CCF database through two complementary perspectives. First, we present key distributional characteristics of the annotated data to show patterns in how climate discourse manifests across 9.2 million sentences spanning temporal, geographic, and thematic dimensions. These baseline distributions provide interesting context for understanding the Canadian climate media landscape and identifying areas for research. Second, we showcase four illustrative analytical applications that leverage the database's granular annotation to uncover patterns invisible to traditional content analysis methods (from the network structure of epistemic authorities to regional variations in

scientific framing and systematic relationships between political actors and editorial prominence). These examples represent only a fraction of the database’s analytical potential; more sophisticated statistical modeling incorporating external data sources and advanced computational methods are currently being developed.<sup>11</sup>

### 3.1 Distribution of Annotation Categories Across the Database

To understand the composition and characteristics of the CCF database, we present a short analysis of how the 65 annotation categories are distributed across our corpus of 9.2 million sentences.

Figure 9 presents the distribution of three key annotation groups across the database. The *Messengers* panel reveals that public officials (33.6%) and natural scientists (7.8%) are the most frequently cited actors in climate coverage, while economic experts (6.5%) and cultural figures (3.2%) receive less attention. This distribution suggests that climate discourse in Canada remains primarily framed through political and scientific lenses. The *Solutions* panel demonstrates an overwhelming emphasis on mitigation strategies (68.7%) compared to adaptation measures (3.1%), indicating that Canadian media coverage focuses predominantly on reducing emissions rather than adaptation. The *Events* panel shows that meetings and conferences (29.0%) and policy announcements (24.9%) dominate event coverage, while protests (0.8%) and cultural events (1.4%) receive minimal attention, suggesting an institutional bias in climate reporting.

---

<sup>11</sup>For instance, we are developing advanced media cascade detection algorithms that leverage the full cross-referenced metadata (journalists, outlets, temporal patterns) to identify information diffusion dynamics across the Canadian media ecosystem. See <https://github.com/antoinelemor/CCF-media-cascade-detection> for ongoing work in this direction.

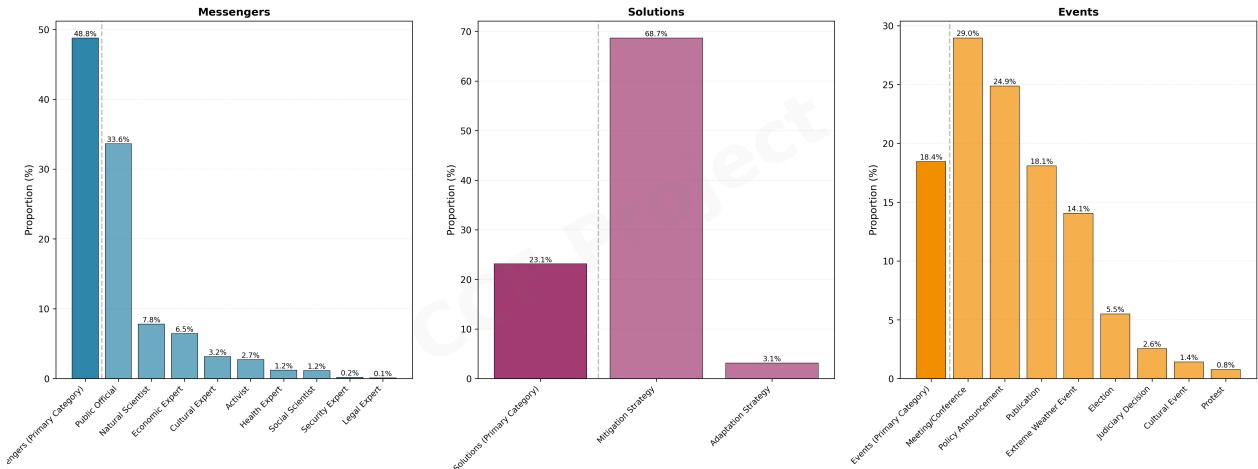


Figure 9: Distribution of annotation categories across the CCF database. The figure shows the average proportion of sentences containing each category for three major annotation groups: *Messengers* (left), *Solutions* (center), and *Events* (right). Primary categories represent sentences with any mention of the respective dimension, while subcategories provide more granular classification (see Table A1 in Appendix for detailed category definitions).

The temporal dynamics presented in Figure 10 reveal great shifts in climate discourse over the last five decades.<sup>12</sup> The most striking pattern is the rise of political framing, which increased from virtually absent in the early 1980s to become the dominant frame by the mid-1990s, stabilizing around 35-40% of coverage in recent years. This politicization corresponds with a dramatic decline in scientific framing, which fell from approximately 40% in the late 1970s to less than 10% by 2025. The economic frame has remained remarkably stable at 15-20% throughout the period, suggesting consistent attention to the financial dimensions of climate change. The environmental frame, focusing on biodiversity and ecosystem impacts, has maintained a steady but modest presence around 3-5% of coverage. These patterns indicate a transformation from climate change as primarily a scientific issue to one dominated by political discourse, with implications for public understanding and policy development.

Figure 11 reveals remarkable consistency in frame distribution across Canada's provinces and suggests a nationally coherent climate discourse despite regional differences in economic structure and climate impacts. Ontario exhibits the highest proportion of political framing (44.4%), possibly reflecting its role as the national political center and most populous province. The Prairie provinces (Alberta 36.9%, Saskatchewan 39.7%) show relatively high political framing combined with stronger economic framing (20.9% and 20.8% respectively), likely reflecting debates over energy transition in these oil-producing regions. Quebec's distribution (Political: 33.3%, Economic: 15.2%, Scientific: 6.6%) aligns closely with the national

<sup>12</sup>All proportions in this analysis represent the average proportion of sentences per article containing each frame, calculated at the article level before temporal or geographic aggregation.

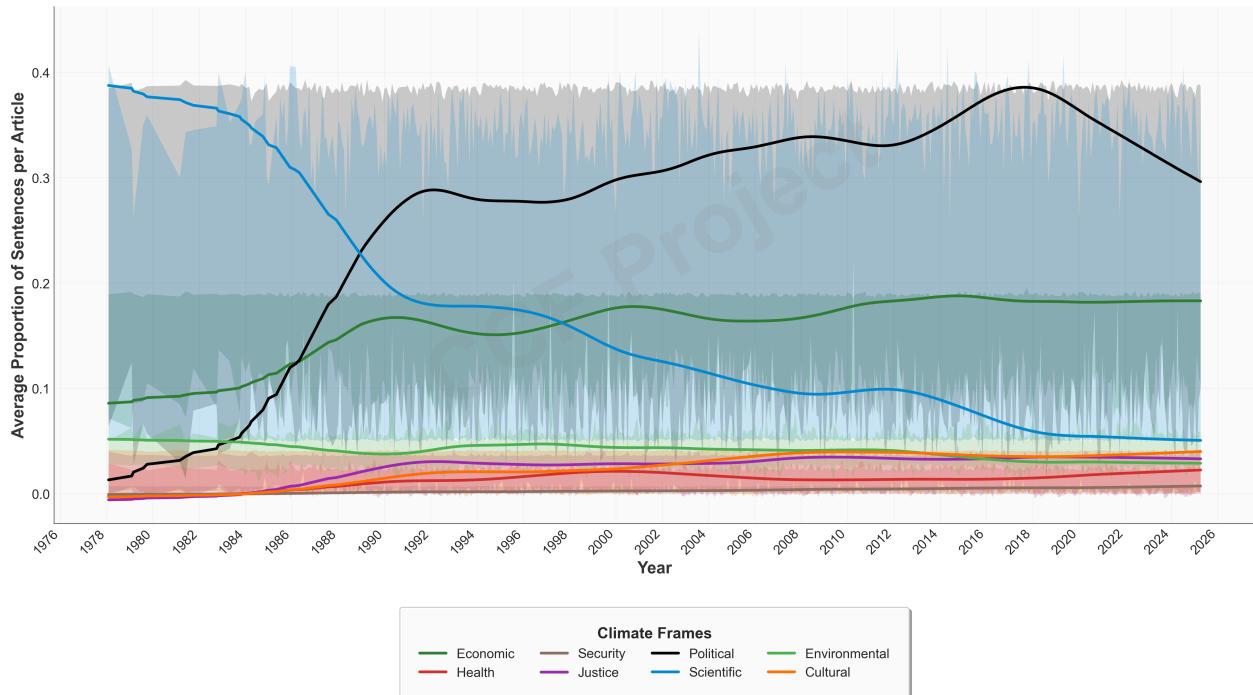


Figure 10: Temporal evolution of climate frames in Canadian media from 1978 to 2025. Lines represent LOESS-smoothed proportions of sentences containing each frame, with shaded areas indicating 95% confidence intervals calculated through bootstrap resampling ( $n=100$ ).

average, while the Yukon shows the highest proportion of scientific framing (10.5%) among all provinces, possibly due to the Arctic's role as a climate change indicator region.

These distributional analyses demonstrate several key characteristics of the CCF database. First, the comprehensive coverage across all 65 annotation categories validates our framework's ability to capture the multidimensional nature of climate discourse. Second, the temporal and geographic patterns reveal how climate communication has evolved from scientific to political domains while maintaining remarkable consistency across provinces. Third, the dominance of institutional actors and events suggests opportunities for diversifying climate narratives to include more grassroots perspectives and adaptation strategies. Together, these patterns provide researchers with a baseline understanding of Canadian climate discourse that can inform more targeted analyses of specific frames, time periods, or regions.

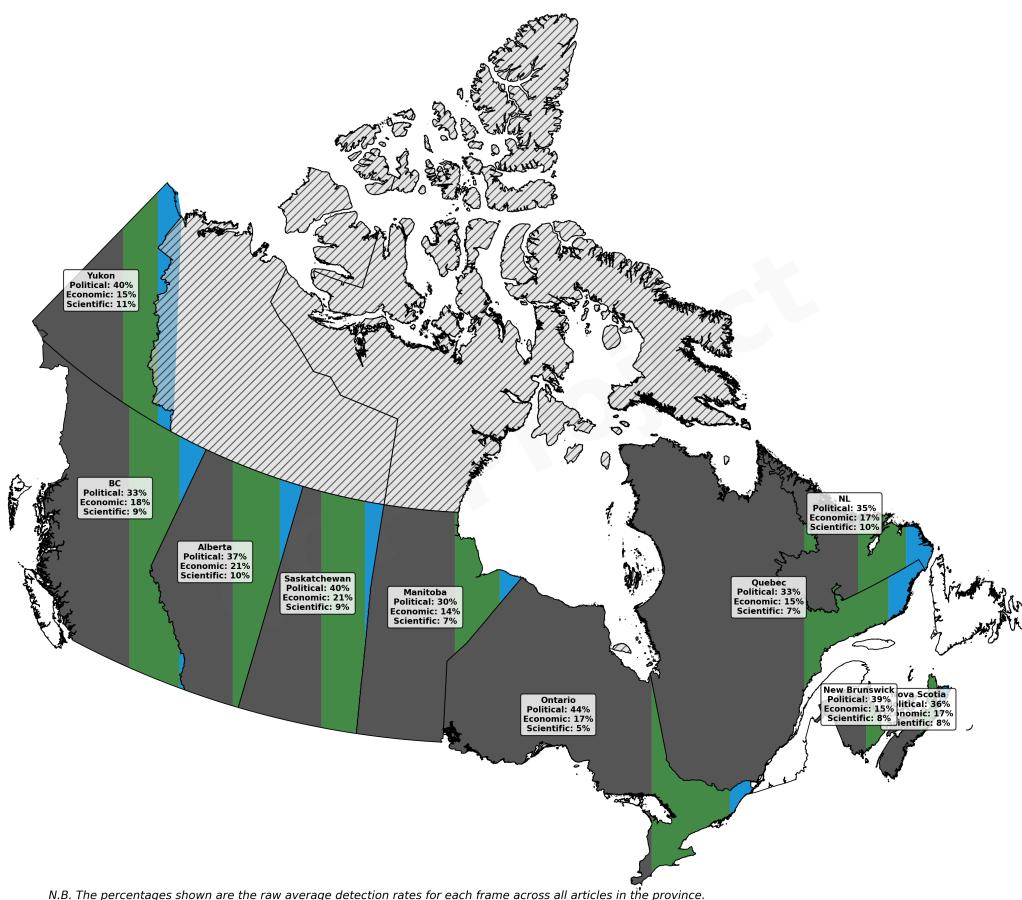


Figure 11: Geographic distribution of the top three climate frames across Canadian provinces. Horizontal bands within each province represent the relative proportions of the three most prevalent frames, with band width corresponding to frame prominence. Provinces shown with hatching indicate no data for analysis.

## 3.2 Examples of Analytical Applications

While the descriptive statistics provide a foundation, the true potential of the CCF database lies in its capacity for sophisticated yet accessible analyses as data are already preprocessed, annotated and validated. Here we present four scientifically grounded applications that showcase this potential: (1) first, we examine how specific political actors systematically achieve higher editorial prominence (i.e., media salience measured as first-page placement probability); (2) second, we map the geographic polarization of climate science discourse across provinces; (3) third, we analyze how different climate frames and themes (e.g., economy, politics, science, environment, justice, culture) influence newspapers' decisions about front-page placement; and (4) fourth, we visualize the network structure of epistemic authorities—identified through our *Messenger/actor* detection categories and named entity recognition, these are the individuals and organizations cited as sources and whose co-citation can reveal the underlying social architecture of climate communication in Canadian media.

### 3.2.1 Political Entities and Editorial Prominence

The CCF database can enable analysis of how specific actors, such as politicians (and others, see Table A1), can influence editorial decisions in climate coverage. As illustrated in Figure 12, we can identify political entities mentioned in at least 50 sentences and calculates their

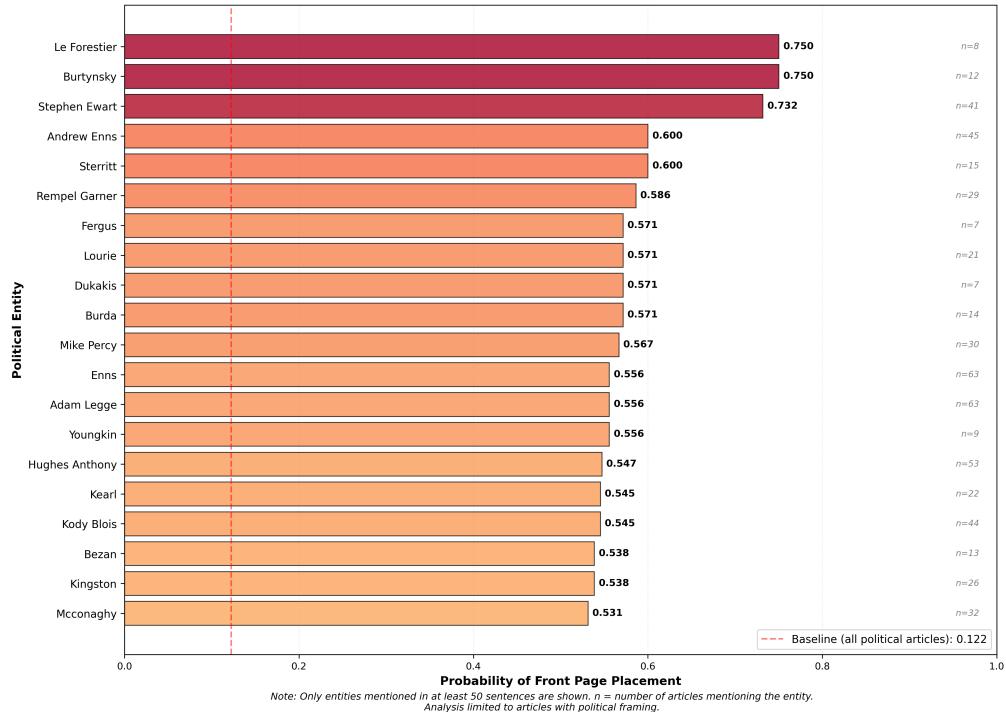


Figure 12: Front page placement probability for political entities mentioned in articles with political framing. Bars represent the probability that an article mentioning each entity appears on the front page, with error bars showing 95% confidence intervals. The baseline (0.122) indicates the average front page probability for all articles with political framing.

associated probability of front page placement. Several patterns emerge from this analysis. First, certain political figures consistently achieve front page placement rates exceeding 60%, suggesting their involvement signals newsworthy climate developments. Second, the variation in placement probability (ranging from 0.531 to 0.750) indicates differential news value assigned to various political actors, potentially reflecting their policy influence or controversial positions. Third, all entities shown exceed the baseline rate of 0.122 for political articles, confirming that named political actors enhance news prominence. This pattern analysis capability enables researchers to study how specific actors shapes editorial decisions and public attention to climate issues.

### 3.2.2 Geographic Polarization of Climate Science Discourse

The CCF database can also be used to produce a detailed geographic analysis of how climate science is discussed across Canadian provinces. Figure 13 reveals interesting regional patterns

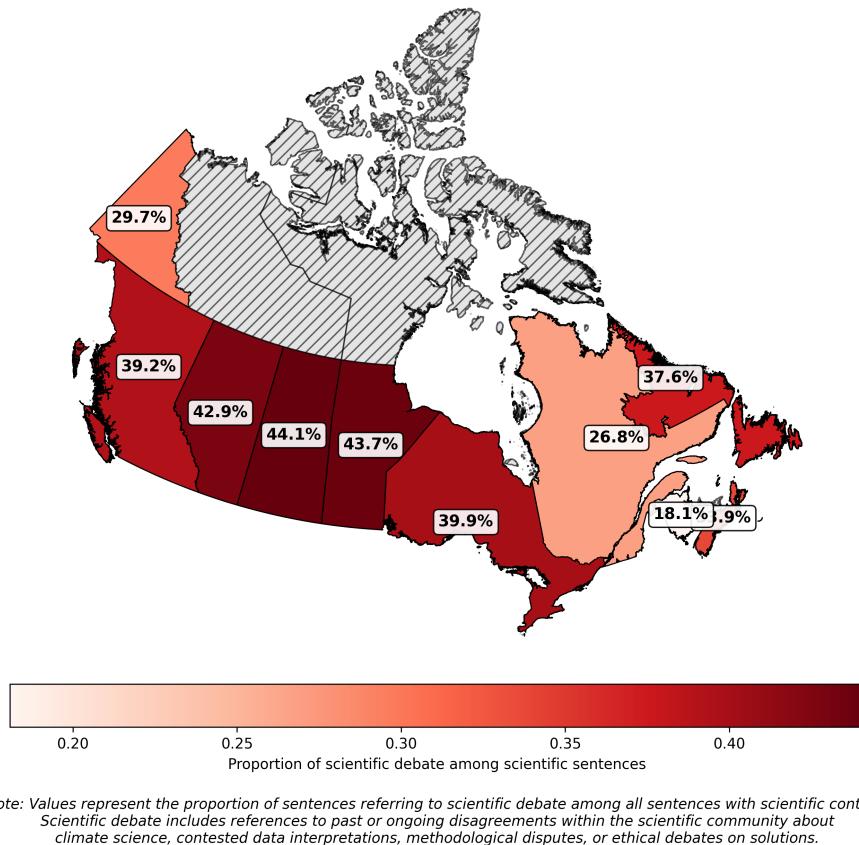


Figure 13: Geographic distribution of scientific debate in Canadian climate coverage. Values represent the proportion of sentences referring to scientific debate among all sentences with scientific content (*Scientific Frame=1*). Scientific debate includes references to past or ongoing disagreements within the scientific community about climate science, contested data interpretations, methodological disputes, or ethical debates on solutions.

in the portrayal of scientific debate. The analysis focuses on our *Scientific debate* category (see Table A1 in Appendix), which captures references to disagreements within the scientific community about climate science, contested interpretations of data, methodological disputes, or ethical debates surrounding proposed solutions.

The analysis shows pronounced geographic variation in how climate science is framed as a matter of debate. Among sentences with scientific content, the proportion presenting science as contested ranges from 18.1% in New Brunswick to 44.1% in Saskatchewan. The Prairie provinces—Saskatchewan (44.1%), Manitoba (43.7%), and Alberta (42.9%)—show remarkably consistent high rates of scientific debate framing which suggests that regions with significant fossil fuel economies are more likely to portray climate science as unsettled. Ontario (39.9%) and British Columbia (39.2%) also exhibit elevated rates, while Quebec (26.8%) and the Atlantic provinces show lower proportions. This pattern indicates that nearly 40% of scientific content in western Canadian media frames climate science through the lens of ongoing debate, despite the overwhelming scientific consensus on anthropogenic climate change.

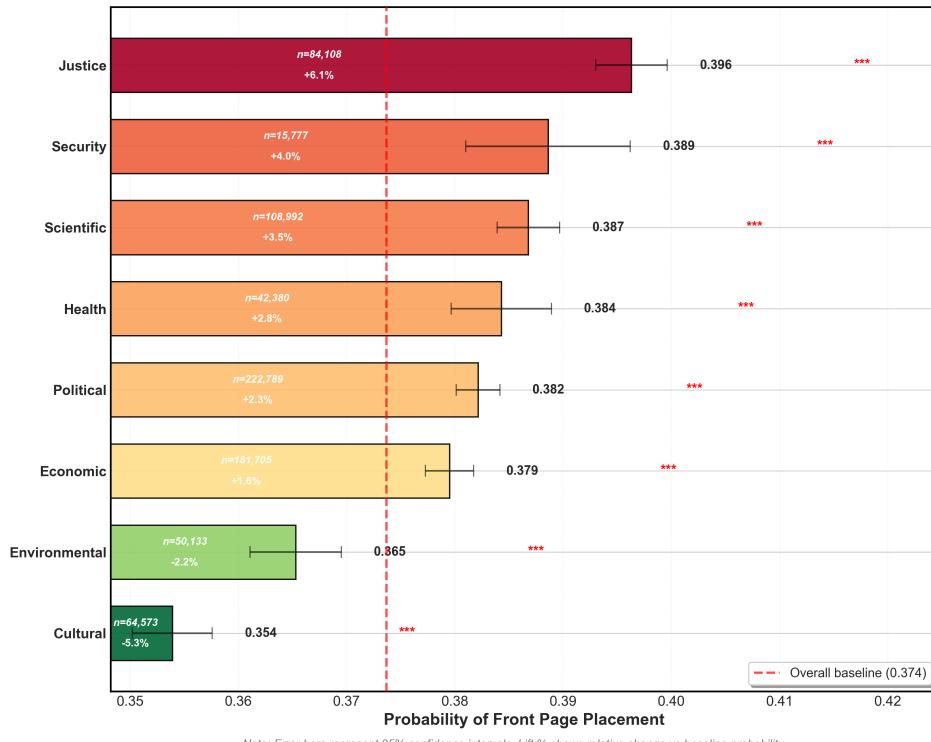


Figure 14: Front page placement probability by climate frame compared to the overall baseline (0.374). Articles containing justice framing achieve the highest placement probability (0.396), representing a 6.1% increase over baseline.

### 3.2.3 Frame-Based Editorial Prioritization

The CCF database can also reveal patterns in how different climate frames influence editorial decisions about story prominence. This analysis examines the relationship between thematic frames and front page placement probability.

Figure 14 shows a clear hierarchy in editorial prioritization: *Justice* (0.396), *Security* (0.389), and *Scientific* (0.387) frames achieve the highest placement rates, all significantly exceeding the baseline of 0.374 ( $p < 0.001$ ). *Political* (0.382) and *Economic* (0.379) frames, despite their prevalence in the database, show more modest increases in placement probability. Notably, *Environment* (0.365) shows a negative association with front page placement, suggesting that biodiversity and ecosystem impacts receive lower editorial priority despite their fundamental importance to climate science. The *Culture* frame demonstrates the strongest

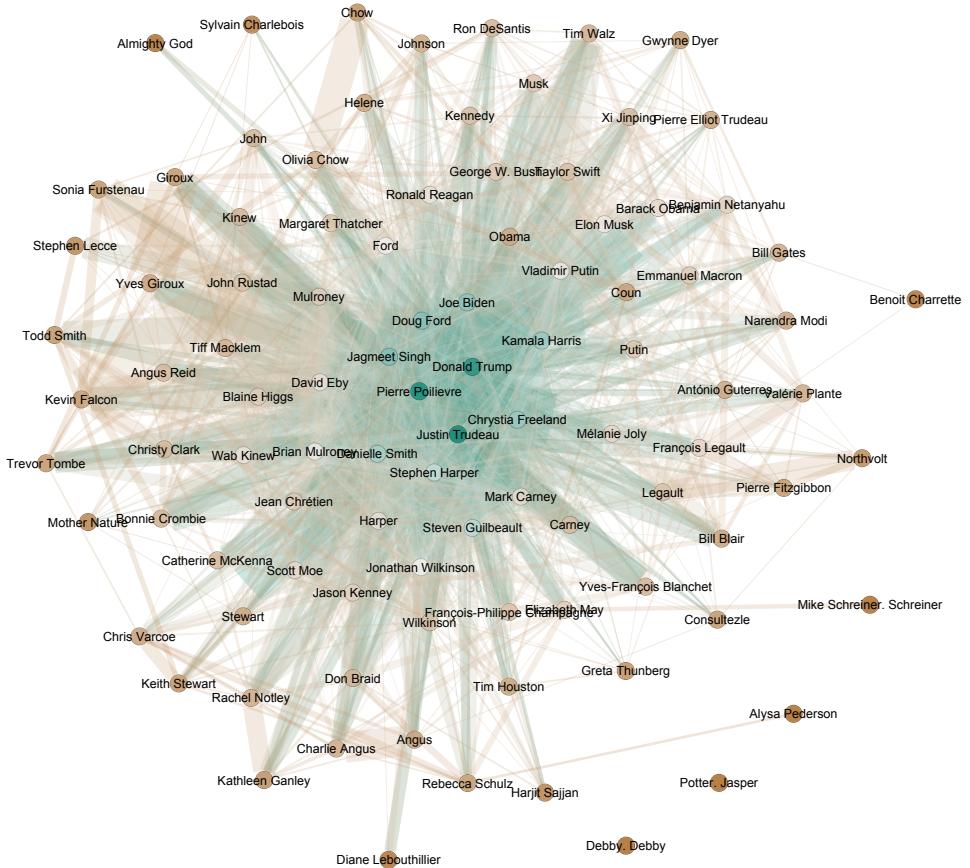


Figure 15: Co-citation network of epistemic authorities in Canadian climate media discourse (year=2024). Node size represents degree centrality, edge thickness indicates co-citation frequency, and colors denote community detection clusters. The network reveals a hierarchical structure with political figures at the core (Trudeau: degree centrality 0.899, Poilievre: 0.879, Trump: 0.848) surrounded by specialized thematic clusters including environmental policy experts (Guilbeault, McKenna), provincial leaders (Ford, Smith, Legault), and economic authorities (Carney, Bank of Canada officials).

negative association (-5.3%). These patterns reveal editorial biases that privilege certain narrative frames over others, potentially shaping public understanding of climate change priorities.

### 3.2.4 Network Structure of Epistemic Authorities

The CCF database can also enable sophisticated network analysis of "epistemic authorities"—i.e., the individuals and organizations cited as sources in climate discourse. By combining our *Messenger* detection categories with Named Entity Recognition, we can systematically identify who speaks about climate change and map their co-citation patterns across journalists and media outlets to reveal the underlying social structure in climate discourse. We could even select specific source by types (e.g., *Expert/Scientist*, *Public Official*, *Activist*, etc.).

Here we present a general analysis of *Messenger* for the year 2024.<sup>13</sup> Figure 15 visualizes the 2024 network topology and shows a highly centralized structure dominated by political figures. Justin Trudeau (2,295 citations), Pierre Poilievre (1,455 citations), and Donald Trump (1,237 citations) form the central core, while scientific authorities occupy peripheral positions. The network exhibits high clustering (coefficient = 0.629), with distinct communities forming around policy domains: environmental ministers cluster together (e.g., Steven Guilbeault co-appears with other environmental officials in 67% of his mentions), provincial premiers form regional governance clusters, and international figures serve as bridges to global discourse while maintaining lower centrality than domestic politicians.

The dominance of political over scientific authorities—with the top 20 authorities including 18 politicians and only 2 scientists—reveals that Canadian media frames climate change primarily through political rather than scientific lenses. Combined with our earlier finding that 40% of science-related content employs debate framing, this network structure suggests that expertise itself becomes politicized in climate communication. The high clustering coefficient indicates that media coverage reinforces echo chambers rather than fostering cross-domain dialogue, with authorities cited within specialized communities that rarely intersect.

## 4 Conclusion

The Canadian Climate Framing (CCF) database establishes a new paradigm for climate communication research through an integration of scale, granularity, and machine learning methodologies. By transforming 266,271 articles spanning 47 years into 9.2 million analysis-ready sentences with 65 validated annotations, we have created not merely a dataset but a comprehensive research infrastructure that aims to deepen how scholars can investigate climate discourse. The achievement of  $F1=0.866$  across all annotation categories demonstrates that computational methods can capture the semantic complexity of climate communication

---

<sup>13</sup>The network analysis employs a multi-stage pipeline: (1) messenger markers are detected through our pre-annotated categories (*Expert/Scientist*, *Public Official*, *Activist*, etc.); (2) NER extracts entities from sentences containing messenger markers; (3) entities are cross-referenced and consolidated using both automated resolution and manual verification for high-profile figures; (4) co-occurrence networks are constructed from article-level co-citations. For visualization clarity, we present the top 100 authorities by citation frequency. For more details, see <https://github.com/antoinelemor/CCF-media-cascade-detection>.

with sufficient accuracy for rigorous academic research. Our methodology’s four-phase architecture provides a reproducible framework that addresses longstanding limitations in climate communication research. The stratified validation using root-inverse probability weighting ensures robust performance even for categories appearing in less than 0.01% of sentences, demonstrating that rare but important discourse patterns need not be sacrificed for computational efficiency.

The CCF database enables fine-grained temporal analysis that was previously impossible to conduct at this scale. Our example of analytical application identified precise inflection points: scientific frames declined from 40% to under 10% between 1990 and 2020. The database’s event detection capabilities (8 sub-categories: extreme weather events, meeting/conference, publication, election, etc.) can allow researchers to correlate these shifts with specific triggers. This granularity transforms broad narratives into testable hypotheses—researchers can now quantify whether political transitions systematically alter discourse, measure the decay rate of scientific authority following policy changes, or test whether economic crises predictably marginalize environmental concerns; among multiple others hypotheses and streams of research.

Geographic analysis can enable researchers to understand how regional economies structure epistemic frameworks. Our analytical application examples show that prairie provinces frame climate science as debate at rates exceeding 40%, while Quebec maintains rates below 20%—a pattern that emerges only when we cross-reference geographic metadata with frame annotations across 9.2 million sentences. This capability allows researchers to test whether fossil fuel dependence correlates with scientific contestation, whether linguistic differences shape climate understanding, or whether provincial policy variations influence media coverage. The database provides the statistical power to distinguish genuine regional differences from sampling artifacts through its comprehensive coverage of 20 newspapers across all provinces.

The combination of messenger detection and Named Entity Recognition also enables systematic mapping of authority and social networks. Our analytical example of co-citation analysis across 266,271 articles shows that politicians comprise 90% of the 20 most-cited authorities, yet more importantly, it quantifies their network positions: political figures maintain greater degree centrality than scientists. This network structure data allows researchers to test whether centrality predicts policy influence, whether clustering patterns reveal discourse coalitions, or whether network position correlates with media framing choices. The database also supports our media cascade detector (<https://github.com/antoinelemor/CCF-media-cascade-detection>), which traces how specific frames propagate across outlets and identifies the originating sources of dominant narratives.

The CCF Database can also shed light on editorial prioritization patterns and reveal systematic biases in news values. Justice and security frames achieve 39.6% and 38.9% front-page placement respectively, while environmental frames achieve only 36.5%—despite environmental frames appearing more frequently overall (12.4% versus 8.7% for justice). This discovery required simultaneous analysis of frame type, page placement metadata, and publication timing across millions of sentences. Researchers can now investigate whether editorial decisions reflect audience preferences, advertiser influence, or journalistic norms, and they can track how these patterns vary across outlets, regions, and time periods.

The database’s distribution characteristics themselves also enable new research possi-

bilities. With 49.68% of sentences containing thematic frames, 48.77% containing messengers, and 18.44% containing events, researchers can analyze frame co-occurrence patterns, messenger-frame associations, and event-driven discourse shifts. The preservation of sentence-level context allows investigation of rhetorical strategies: how frames combine within arguments, how messengers invoke different frames for different audiences, or how emotional tone modulates frame effectiveness. These analyses were previously impossible without manual coding of millions of sentences.

Future developments will expand the CCF framework in two strategic directions. First, we will transform the database into a real-time observatory through continuous data gathering. Daily updates will capture new articles, apply our trained models for immediate annotation, and generate alerts when significant frame shifts occur. This living database will enable researchers to monitor discourse evolution as it unfolds. Second, we would like to extend the framework internationally by adapting our methodology to other national contexts. These parallel databases will enable comparative analysis of how different political economies, media structures, and cultural contexts shape climate discourse, while our cascade detection algorithms will be able to trace how narratives propagate across national boundaries and identify whether certain frames originate in specific countries before spreading globally.

The ultimate significance of the CCF database lies not in its technical achievements but in its potential to advance climate communication research as evidence-informed science. By revealing how climate discourse has evolved, who controls it, and which narratives dominate or disappear, we provide the empirical foundation necessary for designing more effective, inclusive, and urgent climate communication strategies.

## A Appendix: Complete Framework and Performance Metrics

Table A1: Complete CCF annotation framework: All 65 categories with operational definitions

#	Category	Description
<b>THEMATIC FRAMES</b>		
<i>Economic Frame</i>		
1	Economic Frame Detection	Presence of any economic dimension of climate change
2	Negative Economic Impacts	Economic losses from climate change (e.g., floods, infrastructure damage)
3	Positive Economic Impacts	Economic benefits from climate action (e.g., green jobs, energy savings)
4	Costs of Climate Action	Economic costs of transitioning to low-carbon economy
5	Benefits of Climate Action	Economic opportunities from climate investments
6	Economic Sector Footprint	Carbon intensity of economic sectors
<i>Health Frame</i>		
7	Health Frame Detection	Presence of any relationship between climate and health
8	Negative Health Impacts	Heat stress, disease spread, respiratory issues, mental health burdens, mortality
9	Health Co-benefits	Better air quality, improved diets, avoided premature deaths, mental well-being
<i>Security Frame</i>		
10	Security Frame Detection	Presence of any security dimension
11	Climate-Driven Displacement	Military management of evacuations or refugee camps
12	Resource Conflict	Tensions or violence over water, land or minerals worsened by climate change
13	<i>Military Disaster Response*</i>	Army called in for fires, floods, evacuations or relief
14	<i>Military Base Disruption*</i>	Climate threats to military facilities or readiness
<i>Justice Frame</i>		
15	Justice Frame Detection	Presence of any social justice angle
16	Winners & Losers	Groups that benefit or suffer from climate measures (workers, vulnerable populations)
17	North-South Responsibility	Common-but-differentiated responsibilities between high and low-income countries
18	Unequal Impacts	Disproportionate climate impacts on vulnerable groups
19	Unequal Access	Differential access to resources, adaptation, or clean technology
20	Intergenerational Justice	Responsibilities to future generations
<i>Political Frame</i>		
21	Political Frame Detection	Presence of any policy measure or political discussion
22	Policy Measures	Concrete climate laws, regulations, or programmes under debate or in force

Continued on next page

Table A1 – *Continued from previous page*

#	Category	Description
23	Political Debate	Parliamentary disputes, party platforms
24	Political Positioning	Strategic positioning or partisan signalling around climate
25	Public Opinion	Polls or surveys on climate and energy
	<i>Scientific Frame</i>	
26	Scientific Frame Detection	Presence of any scientific aspect
27	Scientific Controversy	Debates on climate change reality, causes, thresholds, geoengineering ethics
28	Discovery & Innovation	New findings on climate impacts or emerging technologies (e.g., carbon capture)
29	Scientific Uncertainty	Expressions of doubt or uncertainty about climate science
30	Scientific Certainty	Strong consensus statements about climate science
	<i>Environmental Frame</i>	
31	Environmental Frame Detection	Presence of any biodiversity concern
32	Habitat Loss	Glacier melt, coral bleaching, forest die-off, wetland drying
33	Species Loss	Local or global extinction risk for animals or plants
	<i>Cultural Frame</i>	
34	Cultural Frame Detection	Presence of any cultural aspect
35	Artistic Representation	Books, documentaries, plays, exhibitions portraying climate themes
36	Event Disruption	Sports or cultural events threatened or cancelled due to climate conditions
37	Loss of Indigenous Practices	Erosion of traditional hunting, fishing, or cultural rituals linked to climate
38	Cultural Sector Footprint	Emissions from film production, fashion, large festivals, etc.
<b>PRIMARY CATEGORIES</b>		
	<i>Actors/Messengers</i>	
39	Actors/Messengers Detection	Presence of any messenger, expert or authority figure
40	Health Expert	Physicians, epidemiologists, health ministers, public health officials
41	Economic Expert	Economists, finance ministers, market analysts, central bank officials
42	Security Expert	Military officers, defense strategists, security scholars
43	Legal Expert	Lawyers, judges, legal scholars, justice ministers
44	Cultural Expert	Artists, writers, athletes, arts scholars commenting on climate change
45	Natural Scientist	Natural science researchers or academics
46	Social Scientist	Social science researchers or academics
47	Activist	Environmental NGO spokespeople or well-known climate activists
48	Public Official	Civil servants or government officials speaking in an official capacity
	<i>Events</i>	

Table A1 – *Continued from previous page*

#	Category	Description
49	Event Detection	Mentions at least one specific event type
50	Extreme Weather Event	Arrival or unfolding of floods, wildfires, hurricanes, heatwaves, etc.
51	Meeting/Conference	International meetings such as COP, UN summits, major national conferences
52	Publication	Publication of governmental, NGO or scientific reports (e.g., IPCC, Lancet)
53	Election	Climate issues raised during local, provincial or national elections
54	Policy Announcement	Debut or unveiling of new climate laws, regulations or action plans
55	Judiciary Decision	Court rulings, trials, regulatory hearings on climate or environment
56	Cultural Event	Sports, artistic, or cultural events (Olympics, marathon, concert, theatre)
57	Protest	Demonstrations or strikes (e.g., climate strike, anti-pipeline protest)
<i>Solutions</i>		
58	Solutions Detection	Mentions any mitigation or adaptation measure
59	Mitigation Strategy	Measures to reduce GHG emissions or enhance carbon sinks
60	Adaptation Strategy	Measures to increase social or ecological resilience to climate impacts
<b>EMOTIONAL TONE</b>		
61	Positive Emotion	Hope, optimism, enthusiasm about climate solutions
62	Negative Emotion	Fear, anxiety, despair about climate impacts
63	Neutral Emotion	No clear emotional tone, factual reporting
<b>GEOGRAPHIC FOCUS</b>		
64	Canadian Context	Canadian places, actors, data, and policies
<b>URGENCY/ALARMISM</b>		
65	Urgency/Alarmism	Conveys immediate danger, crisis, or “code red” urgency

\* Insufficient training data (fewer than 10 positive examples in training set).

Table A3: Complete model training performance metrics for all annotation categories

#	Category	F1 (Class 1)		F1 (Class 0)		Macro F1		
		EN	FR	EN	FR	EN	FR	
<b>THEMATIC FRAMES</b>								
<i>Economic Frame</i>								
1	Economic Frame Detection	0.745	0.814	0.944	0.957	0.845	0.885	
2	Negative Economic Impacts	0.727	0.889	0.914	0.944	0.821	0.917	
3	Positive Economic Impacts	0.333	0.400	0.995	0.996	0.664	0.698	
4	Costs of Climate Action	0.615	0.500	0.848	0.913	0.732	0.707	
5	Benefits of Climate Action	0.500	0.516	0.952	0.981	0.726	0.749	
6	Economic Sector Footprint	0.857	0.857	0.938	0.950	0.897	0.904	
<i>Health Frame</i>								
7	Health Frame Detection	0.800	0.667	0.989	0.994	0.894	0.830	
8	Negative Health Impacts	0.909	0.857	0.000	0.000	0.455	0.429	
9	Positive Health Impacts	—	—	—	—	—	—	
10	Health Co-benefits	0.400	0.011	0.571	0.227	0.486	0.119	
11	Health Sector Footprint	—	—	—	—	—	—	
<i>Security Frame</i>								
12	Security Frame Detection	0.870	0.898	0.996	0.997	0.933	0.898	
13	Military Disaster Response	—	0.026	—	1.000	—	0.026	
14	Military Base Disruption	—	1.000	—	1.000	—	1.000	
15	Climate-Driven Displacement	1.000	1.000	1.000	1.000	1.000	1.000	
16	Resource Conflict	1.000	1.000	1.000	1.000	1.000	1.000	
17	Defense Sector Footprint	—	—	—	—	—	—	
<i>Justice Frame</i>								
18	Justice Frame Detection	0.719	0.717	0.975	0.981	0.847	0.849	
19	Winners & Losers	0.667	0.667	0.933	0.923	0.800	0.795	
20	North-South Responsibility	0.857	0.750	0.909	0.750	0.883	0.750	
21	Unequal Impacts	0.571	0.600	0.992	0.995	0.782	0.798	
22	Unequal Access	0.364	0.625	0.000	0.993	0.182	0.809	
23	Intergenerational Justice	0.933	0.800	0.999	0.999	0.966	0.899	
<i>Political Frame</i>								
24	Political Frame Detection	0.808	0.774	0.897	0.888	0.853	0.831	
25	Policy Measures	0.621	0.667	0.985	0.966	0.803	0.816	

Continued on next page

Table A3 – *Continued from previous page*

#	Category	F1 (Class 1)		F1 (Class 0)		Macro F1	
		EN	FR	EN	FR	EN	FR
26	Political Debate	0.916	0.966	0.222	0.571	0.569	0.768
27	Political Positioning	0.750	0.800	0.946	0.989	0.848	0.895
28	Public Opinion	0.909	1.000	0.999	1.000	0.954	1.000
	<i>Scientific Frame</i>						
29	Scientific Frame Detection	0.784	0.702	0.953	0.962	0.869	0.832
30	Scientific Controversy	0.737	0.500	0.848	0.875	0.793	0.688
31	Discovery & Innovation	0.828	0.938	0.800	0.750	0.814	0.844
32	Scientific Uncertainty	0.727	0.333	0.927	0.995	0.827	0.664
33	Scientific Certainty	0.571	0.000	0.933	0.919	0.752	0.459
	<i>Environmental Frame</i>						
34	Environmental Frame Detection	0.842	0.625	0.989	0.980	0.915	0.802
35	Habitat Loss	1.000	0.857	1.000	0.800	1.000	0.829
36	Species Loss	0.889	0.857	0.857	0.800	0.873	0.829
	<i>Cultural Frame</i>						
37	Cultural Frame Detection	0.773	0.833	0.986	0.993	0.879	0.913
38	Artistic Representation	1.000	1.000	1.000	1.000	1.000	1.000
39	Event Disruption	0.706	1.000	0.993	1.000	0.850	1.000
40	Loss of Indigenous Practices	1.000	0.026	1.000	0.899	1.000	0.462
41	Cultural Sector Footprint	1.000	0.005	1.000	0.143	1.000	0.074
	<b>PRIMARY CATEGORIES</b>						
	<i>Actors/Messengers</i>						
42	Actors/Messengers Detection	0.912	0.929	0.904	0.915	0.908	0.922
43	Health Expert	0.857	0.909	0.997	0.999	0.927	0.954
44	Economic Expert	0.600	0.750	0.970	0.973	0.785	0.861
45	Security Expert	0.667	1.000	0.993	1.000	0.830	1.000
46	Legal Expert	0.667	0.769	0.999	0.996	0.833	0.883
47	Cultural Expert	1.000	0.556	1.000	0.990	1.000	0.773
48	Natural Scientist	0.789	0.833	0.977	0.971	0.883	0.902
49	Social Scientist	0.769	0.645	0.996	0.986	0.883	0.816
50	Activist	0.571	1.000	0.978	1.000	0.775	1.000
51	Public Official	0.703	0.923	0.897	0.964	0.800	0.943
	<i>Events</i>						
52	Event Detection	0.794	0.819	0.935	0.932	0.865	0.876

*Continued on next page*

Table A3 – *Continued from previous page*

#	Category	F1 (Class 1)		F1 (Class 0)		Macro F1	
		EN	FR	EN	FR	EN	FR
53	Extreme Weather Event	1.000	0.857	1.000	0.968	1.000	0.912
54	Meeting/Conference	0.824	0.957	0.936	0.982	0.880	0.969
55	Publication	0.873	1.000	0.990	1.000	0.931	1.000
56	Election	1.000	0.800	1.000	0.998	1.000	0.899
57	Policy Announcement	0.769	0.727	0.943	0.954	0.856	0.841
58	Judiciary Decision	1.000	1.000	1.000	1.000	1.000	1.000
59	Cultural Event	0.333	1.000	0.995	1.000	0.664	1.000
60	Protest	0.889	0.727	0.999	0.996	0.944	0.862
<i>Solutions</i>							
61	Solutions Detection	0.737	0.878	0.914	0.944	0.825	0.911
62	Mitigation Strategy	0.750	0.812	0.935	0.942	0.842	0.877
63	Adaptation Strategy	0.696	0.800	0.991	0.989	0.843	0.894
<b>EMOTIONAL TONE</b>							
62	Positive Emotion	0.526	0.690	0.966	0.968	0.746	0.829
63	Negative Emotion	0.706	0.765	0.785	0.843	0.745	0.804
64	Neutral Emotion	0.741	0.789	0.676	0.693	0.709	0.741
<b>GEOGRAPHIC FOCUS</b>							
65	Canadian Context	0.942	0.964	0.968	0.980	0.955	0.972
<b>URGENCY/ALARMISM</b>							
66	Urgency/Alarmism	0.591	0.649	0.975	0.984	0.783	0.816
<b>AVERAGE PERFORMANCE METRICS</b>							
English Average		0.769	–	0.905	–	0.837	–
French Average		–	0.739	–	0.909	–	0.816
Overall Average		0.754		0.907		0.826	
<b>TOTAL: 65 ANNOTATION CATEGORIES</b>							
(62 categories with at least one model; 3 categories entirely excluded*)							

\* Insufficient training data (fewer than 10 positive examples in training set).

Table A5: Training and validation dataset distribution across all annotation categories

#	Category	English				French				
		Training		Validation		Training		Validation		
		Pos	Neg	Pos	Neg	Pos	Neg	Pos	Neg	
<b>THEMATIC FRAMES</b>										
<i>Economic Frame</i>										
1	Economic Frame Detection	218	1067	24	118	252	1164	28	129	
2	Negative Economic Impacts	61	158	6	17	75	178	8	19	
3	Positive Economic Impacts	14	205	1	22	7	245	1	27	
4	Costs of Climate Action	63	156	6	17	42	211	4	23	
5	Benefits of Climate Action	36	183	3	20	48	205	5	22	
6	Economic Sector Footprint	67	152	7	16	77	176	8	19	
<i>Health Frame</i>										
7	Health Frame Detection	57	1228	6	136	37	1379	4	153	
8	Negative Health Impacts	47	10	5	1	32	5	3	1	
9	Positive Health Impacts	0	56	1	6	0	36	1	4	
10	Health Co-benefits	8	49	1	5	4	33	1	3	
11	Health Sector Footprint	0	56	1	6	0	37	0	4	
<i>Security Frame</i>										
12	Security Frame Detection	19	1266	2	140	19	1397	2	155	
13	Military Disaster Response	0	18	1	2	4	15	1	1	
14	Military Base Disruption	0	18	1	2	1	18	1	1	
15	Climate-Driven Displacement	10	9	1	1	7	12	1	1	
16	Resource Conflict	6	13	1	1	7	12	1	1	
17	Defense Sector Footprint	0	19	0	2	0	19	0	2	
<i>Justice Frame</i>										
18	Justice Frame Detection	90	1196	9	132	81	1335	9	148	
19	Winners & Losers	20	70	2	7	20	62	2	6	
20	North-South Responsibility	38	52	4	5	41	41	4	4	
21	Unequal Impacts	8	81	1	9	9	72	1	8	
22	Unequal Access	21	69	2	7	20	62	2	6	
23	Intergenerational Justice	14	76	1	8	6	75	1	8	
<i>Political Frame</i>										
24	Political Frame Detection	416	869	46	96	440	977	48	108	

*Continued on next page*

Table A5 – *Continued from previous page*

#	Category	English				French			
		Training		Validation		Training		Validation	
		Pos	Neg	Pos	Neg	Pos	Neg	Pos	Neg
25	Policy Measures	62	355	6	39	58	382	6	42
26	Political Debate	344	72	38	8	396	45	43	4
27	Political Positioning	61	356	6	39	35	405	3	45
28	Public Opinion	21	396	2	43	20	420	2	46
<i>Scientific Frame</i>									
29	Scientific Frame Detection	243	1042	27	115	185	1232	20	136
30	Scientific Controversy	98	146	10	16	46	139	5	15
31	Discovery & Innovation	126	117	14	13	144	41	16	4
32	Scientific Uncertainty	50	194	5	21	17	169	1	18
33	Scientific Certainty	37	207	4	22	32	153	3	17
<i>Environmental Frame</i>									
34	Environmental Frame Detection	84	1201	9	133	63	1354	6	150
35	Habitat Loss	50	35	5	3	35	28	3	3
36	Species Loss	42	43	4	4	32	31	3	3
<i>Cultural Frame</i>									
37	Cultural Frame Detection	44	1242	4	137	45	1371	5	152
38	Artistic Representation	20	24	2	2	18	27	2	3
39	Event Disruption	11	33	1	3	21	25	2	2
40	Loss of Indigenous Practices	10	34	1	3	4	41	1	4
41	Cultural Sector Footprint	1	42	1	4	2	43	1	4
<b>PRIMARY CATEGORIES</b>									
<i>Actors/Messengers</i>									
42	Actors/Messengers Detection	657	629	72	69	736	681	81	75
43	Health Expert	10	647	1	71	8	728	1	80
44	Economic Expert	55	602	6	66	67	669	7	74
45	Security Expert	3	653	1	72	10	726	1	80
46	Legal Expert	1	655	1	72	4	731	1	81
47	Cultural Expert	13	644	1	71	9	727	1	80
48	Natural Scientist	107	550	11	61	117	620	12	68
49	Social Scientist	9	648	1	71	26	711	2	78
50	Activist	53	604	5	67	55	681	6	75
51	Public Official	171	486	18	54	223	513	24	57

*Continued on next page*

Table A5 – *Continued from previous page*

#	Category	English				French			
		Training		Validation		Training		Validation	
		Pos	Neg	Pos	Neg	Pos	Neg	Pos	Neg
<i>Events</i>									
52	Event Detection	298	987	33	109	355	1062	39	117
53	Extreme Weather Event	72	226	8	25	59	297	6	32
54	Meeting/Conference	75	224	8	24	100	255	11	28
55	Publication	74	225	8	24	114	242	12	26
56	Election	16	283	1	31	18	337	2	37
57	Policy Announcement	55	243	6	27	60	296	6	32
58	Judiciary Decision	10	288	1	32	10	345	1	38
59	Cultural Event	2	296	1	32	11	344	1	38
60	Protest	11	288	1	31	8	347	1	38
<i>Solutions</i>									
61	Solutions Detection	314	972	34	107	415	1001	46	111
<b>EMOTIONAL TONE</b>									
62	Positive Emotion	103	1182	11	131	151	1266	16	140
63	Negative Emotion	491	794	54	88	537	880	59	97
64	Neutral Emotion	686	599	76	66	726	691	80	76
<b>GEOGRAPHIC FOCUS</b>									
65	Canadian Context	445	840	49	93	493	924	54	102
<b>URGENCY/ALARMISM</b>									
66	Urgency/Alarmism	66	1219	7	135	63	1354	6	150

33

Note: Categories with insufficient positive training samples (marked with \* in Table A3) were excluded from final model training and are not shown in this distribution table.

Table A6: Named Entity Recognition model performance metrics

<b>Language</b>	<b>Model</b>	<b>PER F1</b>	<b>ORG F1</b>	<b>LOC F1</b>
English	BERT-base-NER	0.961	0.811	0.925
French	spaCy fr_core_news_lg	0.880	—	—
	CamemBERT-NER	—	0.824	0.929

Note: The hybrid approach for French combines spaCy for person entities (PER) with CamemBERT-NER for organization (ORG) and location (LOC) entities based on empirical evaluation on our dataset.

Performance metrics are from the original model documentation.

Table A8: Detailed validation performance metrics for trained models (62 categories)

#	Category	F1 Macro			F1 Micro			F1 Weighted			Support			
		EN	FR	ALL	EN	FR	ALL	EN	FR	ALL	EN	FR	ALL	
<b>THEMATIC FRAMES</b>														
<i>Economic Frame</i>														
1	Economic Frame Detection	0.847	0.814	0.830	0.892	0.864	0.878	0.890	0.862	0.875	120	129	249	
2	Negative Economic Impacts	0.820	0.741	0.783	0.848	0.810	0.828	0.856	0.833	0.842	26	18	44	
3	Positive Economic Impacts	0.806	0.508	0.684	0.905	0.845	0.873	0.913	0.908	0.900	12	1	13	
4	Costs of Climate Action	0.741	0.702	0.723	0.800	0.802	0.801	0.813	0.806	0.810	22	23	45	
5	Benefits of Climate Action	0.885	0.817	0.848	0.924	0.871	0.896	0.928	0.872	0.899	19	26	45	
6	Economic Sector Footprint	0.785	0.788	0.787	0.838	0.845	0.842	0.831	0.843	0.837	30	29	59	
<i>Health Frame</i>														
7	Health Frame Detection	0.800	0.817	0.808	0.947	0.961	0.954	0.947	0.964	0.956	35	23	58	
8	Negative Health Impacts	0.386	0.340	0.364	0.629	0.514	0.571	0.485	0.349	0.416	22	18	40	
9	Health Co-benefits	0.252	0.103	0.186	0.257	0.114	0.186	0.301	0.023	0.177	4	4	8	
<i>Security Frame</i>														
10	Security Frame Detection	0.825	0.730	0.782	0.951	0.939	0.945	0.955	0.950	0.952	30	19	49	
11	Climate-Driven Displacement	0.476	0.417	0.481	0.545	0.429	0.488	0.468	0.476	0.459	21	6	27	
12	Resource Conflict	0.358	0.286	0.324	0.364	0.286	0.326	0.398	0.286	0.346	7	6	13	
13	Military Disaster Response	–	0.045	0.045	–	0.048	0.048	–	0.004	0.004	–	2	2	
14	Military Base Disruption	–	1.000	1.000	–	1.000	1.000	–	1.000	1.000	–	2	2	
<i>Justice Frame</i>														
15	Justice Frame Detection	0.832	0.880	0.857	0.917	0.937	0.927	0.921	0.939	0.930	62	74	136	
16	Winners & Losers	0.654	0.737	0.708	0.765	0.762	0.764	0.802	0.775	0.786	10	22	32	
17	North-South Responsibility	0.900	0.823	0.860	0.926	0.833	0.879	0.927	0.838	0.883	19	27	46	
18	Unequal Impacts	0.786	0.894	0.845	0.852	0.917	0.885	0.858	0.918	0.888	16	21	37	
19	Unequal Access	0.475	0.796	0.651	0.481	0.833	0.661	0.456	0.835	0.674	27	23	50	
20	Intergenerational Justice	0.924	0.952	0.942	0.963	0.964	0.964	0.965	0.965	0.965	10	19	29	
<i>Political Frame</i>														
21	Political Frame Detection	0.812	0.842	0.829	0.827	0.844	0.836	0.829	0.845	0.837	167	217	384	
22	Policy Measures	0.802	0.722	0.756	0.914	0.864	0.886	0.909	0.854	0.879	26	38	64	
23	Political Debate	0.469	0.650	0.564	0.694	0.820	0.763	0.595	0.781	0.699	127	175	302	

Continued on next page

Table A8 – *Continued from previous page*

#	Category	F1 Macro			F1 Micro			F1 Weighted			Support		
		EN	FR	ALL	EN	FR	ALL	EN	FR	ALL	EN	FR	ALL
24	Political Positioning	0.752	0.808	0.782	0.801	0.912	0.862	0.821	0.920	0.877	35	24	59
25	Public Opinion	0.877	0.910	0.894	0.962	0.974	0.969	0.966	0.976	0.971	12	15	27
<i>Scientific Frame</i>													
26	Scientific Frame Detection	0.891	0.890	0.890	0.951	0.949	0.950	0.948	0.947	0.947	74	75	149
27	Scientific Controversy	0.962	0.867	0.911	0.962	0.869	0.912	0.962	0.869	0.912	26	26	52
28	Discovery & Innovation	0.884	0.797	0.838	0.885	0.803	0.841	0.886	0.803	0.842	31	36	67
29	Scientific Uncertainty	0.783	0.768	0.775	0.846	0.836	0.841	0.860	0.857	0.858	9	9	18
30	Scientific Certainty	0.769	0.460	0.682	0.846	0.852	0.850	0.862	0.785	0.843	8	9	17
<i>Environmental Frame</i>													
31	Environmental Frame Detection	0.866	0.875	0.871	0.967	0.967	0.967	0.969	0.968	0.969	28	32	60
32	Habitat Loss	0.804	0.573	0.721	0.806	0.707	0.753	0.802	0.637	0.730	16	26	42
33	Species Loss	0.714	0.692	0.703	0.722	0.707	0.714	0.717	0.694	0.705	19	21	40
<i>Cultural Frame</i>													
34	Cultural Frame Detection	0.802	0.581	0.708	0.923	0.880	0.901	0.931	0.920	0.921	41	11	52
35	Artistic Representation	0.834	0.469	0.659	0.851	0.559	0.704	0.858	0.649	0.739	18	6	24
36	Event Disruption	0.631	0.354	0.484	0.731	0.456	0.593	0.777	0.595	0.686	8	2	10
37	Loss of Indigenous Practices	0.824	0.415	0.595	0.910	0.618	0.763	0.922	0.749	0.825	7	1	8
38	Cultural Sector Footprint	0.702	0.014	0.351	0.925	0.015	0.467	0.945	0.014	0.613	2	1	3
<b>PRIMARY CATEGORIES</b>													
<i>Actors/Messengers</i>													
39	Actors/Messengers Detection	0.955	0.967	0.961	0.957	0.969	0.963	0.957	0.969	0.963	301	312	613
40	Health Expert	0.862	0.951	0.907	0.964	0.987	0.975	0.966	0.987	0.976	19	21	40
41	Economic Expert	0.847	0.895	0.874	0.954	0.961	0.957	0.949	0.963	0.957	30	28	58
42	Security Expert	0.728	0.813	0.787	0.970	0.958	0.964	0.973	0.962	0.968	7	14	21
43	Legal Expert	0.861	0.927	0.914	0.990	0.980	0.985	0.989	0.982	0.986	7	19	26
44	Cultural Expert	0.853	0.825	0.843	0.954	0.967	0.961	0.954	0.969	0.961	26	13	39
45	Natural Scientist	0.920	0.896	0.907	0.960	0.941	0.951	0.960	0.944	0.952	44	46	90
46	Social Scientist	0.938	0.948	0.944	0.987	0.984	0.985	0.987	0.984	0.986	17	24	41
47	Activist	0.934	0.903	0.918	0.977	0.964	0.970	0.976	0.963	0.970	31	33	64
48	Public Official	0.927	0.930	0.929	0.934	0.935	0.934	0.934	0.935	0.934	104	113	217
<i>Events</i>													

Continued on next page

Table A8 – *Continued from previous page*

#	Category	F1 Macro			F1 Micro			F1 Weighted			Support		
		EN	FR	ALL	EN	FR	ALL	EN	FR	ALL	EN	FR	ALL
49	Event Detection	0.881	0.912	0.897	0.886	0.913	0.900	0.885	0.913	0.900	201	222	423
50	Extreme Weather Event	0.887	0.891	0.889	0.935	0.955	0.946	0.938	0.958	0.949	29	21	50
51	Meeting/Conference	0.855	0.891	0.877	0.908	0.914	0.911	0.914	0.918	0.916	30	50	80
52	Publication	0.950	0.962	0.958	0.973	0.968	0.970	0.973	0.969	0.971	30	62	92
53	Election	0.911	0.895	0.903	0.962	0.964	0.963	0.960	0.963	0.962	25	22	47
54	Policy Announcement	0.821	0.837	0.829	0.886	0.914	0.901	0.895	0.915	0.906	29	33	62
55	Judiciary Decision	0.628	0.939	0.808	0.903	0.982	0.946	0.927	0.981	0.951	6	19	25
56	Cultural Event	0.906	0.857	0.881	0.973	0.964	0.968	0.971	0.967	0.968	17	12	29
57	Protest	0.982	0.909	0.943	0.995	0.973	0.983	0.995	0.973	0.983	16	18	34
<i>Solutions</i>													
58	Solutions Detection	0.891	0.909	0.900	0.921	0.931	0.926	0.922	0.930	0.926	114	133	247
59	Mitigation Strategy	0.872	0.801	0.839	0.884	0.839	0.861	0.882	0.850	0.864	76	97	173
60	Adaptation Strategy	0.889	0.903	0.895	0.934	0.952	0.943	0.934	0.954	0.944	22	16	38
<b>EMOTIONAL TONE</b>													
61	Positive Emotion	0.712	0.692	0.702	0.902	0.892	0.897	0.910	0.904	0.907	38	37	75
62	Negative Emotion	0.793	0.804	0.799	0.807	0.813	0.810	0.814	0.819	0.816	151	165	316
63	Neutral Emotion	0.756	0.798	0.777	0.764	0.809	0.787	0.765	0.807	0.787	295	306	601
<b>GEOGRAPHIC FOCUS</b>													
64	Canadian Context	0.986	0.978	0.982	0.986	0.978	0.982	0.986	0.978	0.982	235	255	490
<b>URGENCY/ALARMISM</b>													
65	Urgency/Alarmism	0.874	0.835	0.853	0.976	0.965	0.970	0.977	0.967	0.972	23	25	48
<b>OVERALL PERFORMANCE</b>													
All Categories		0.869	0.864	0.866	0.909	0.902	0.905	0.911	0.905	0.908	3,069	3,332	6,401

Note: Metrics represent macro-averaged scores across positive and negative classes for each category. Support indicates the number of positive examples in the validation set (500 sentences per language, 1,000 total). EN = English, FR = French, ALL = Combined. Categories are organized by their hierarchical grouping as implemented in the annotation pipeline.

Table A10: Database-wide distribution of annotation categories across 9.2 million sentences

#	Category	Count			Proportion (%)			
		EN	FR	ALL	EN	FR	ALL	
<b>THEMATIC FRAMES</b>								
<i>Economic Frame</i>								
1	Economic Frame Detection	1,235,686	185,032	1,420,718	15.95	12.64	15.43	
2	Costs of Climate Action	517,576	40,259	557,835	6.68	2.75	6.06	
3	Economic Sector Footprint	293,707	61,974	355,681	3.79	4.23	3.86	
4	Negative Economic Impacts	193,399	30,670	224,069	2.50	2.10	2.43	
5	Benefits of Climate Action	75,339	30,266	105,605	0.97	2.07	1.15	
6	Positive Economic Impacts	1,157	1,949	3,106	0.01	0.13	0.03	
<i>Health Frame</i>								
7	Health Frame Detection	106,777	21,164	127,941	1.38	1.45	1.39	
8	Negative Health Impacts	106,777	21,164	127,941	1.38	1.45	1.39	
9	Health Co-benefits	63,844	19,595	83,439	0.82	1.34	0.91	
<i>Security Frame</i>								
10	Security Frame Detection	32,973	10,255	43,228	0.43	0.70	0.47	
11	Resource Conflict	20,434	9,015	29,449	0.26	0.62	0.32	
12	Climate-Driven Displacement	21,417	4,248	25,665	0.28	0.29	0.28	
13	Military Disaster Response	0	10,253	10,253	0.00	0.70	0.11	
14	Military Base Disruption	0	4	4	0.00	0.00	0.00	
<i>Justice Frame</i>								
15	Justice Frame Detection	277,005	20,938	297,943	3.58	1.43	3.23	
16	Unequal Access	215,729	4,791	220,520	2.78	0.33	2.39	
17	North-South Responsibility	59,851	9,084	68,935	0.77	0.62	0.75	
18	Winners & Losers	45,568	8,905	54,473	0.59	0.61	0.59	
19	Intergenerational Justice	26,810	1,160	27,970	0.35	0.08	0.30	
20	Unequal Impacts	13,216	2,188	15,404	0.17	0.15	0.17	
<i>Political Frame</i>								
21	Political Frame Detection	2,411,443	409,831	2,821,274	31.13	28.00	30.63	
22	Political Debate	2,360,465	401,652	2,762,117	30.47	27.44	29.99	
23	Political Positioning	860,760	49,748	910,508	11.11	3.40	9.89	

*Continued on next page*

Table A10 – *Continued from previous page*

#	Category	Count			Proportion (%)		
		EN	FR	ALL	EN	FR	ALL
24	Policy Measures	205,010	18,833	223,843	2.65	1.29	2.43
25	Public Opinion	46,850	8,186	55,036	0.60	0.56	0.60
<i>Scientific Frame</i>							
26	Scientific Frame Detection	517,413	75,144	592,557	6.68	5.13	6.43
27	Discovery & Innovation	345,362	62,495	407,857	4.46	4.27	4.43
28	Scientific Controversy	212,955	14,929	227,884	2.75	1.02	2.47
29	Scientific Uncertainty	64,848	2,065	66,913	0.84	0.14	0.73
30	Scientific Certainty	48,588	0	48,588	0.63	0.00	0.53
<i>Environmental Frame</i>							
31	Environmental Frame Detection	206,579	74,554	281,133	2.67	5.09	3.05
32	Habitat Loss	107,529	62,541	170,070	1.39	4.27	1.85
33	Species Loss	118,604	49,797	168,401	1.53	3.40	1.83
<i>Cultural Frame</i>							
34	Cultural Frame Detection	290,549	82,465	373,014	3.75	5.63	4.05
35	Artistic Representation	202,492	44,568	247,060	2.61	3.05	2.68
36	Cultural Sector Footprint	13	82,382	82,395	0.00	5.63	0.89
37	Event Disruption	23,009	34,855	57,864	0.30	2.38	0.63
38	Loss of Indigenous Practices	7,201	18,778	25,979	0.09	1.28	0.28
<b>PRIMARY CATEGORIES</b>							
<i>Actors/Messengers</i>							
39	Actors/Messengers Detection	3,848,262	644,020	4,492,282	49.68	44.00	48.77
40	Public Official	1,286,959	223,649	1,510,608	16.61	15.28	16.40
41	Natural Scientist	279,864	70,656	350,520	3.61	4.83	3.81
42	Economic Expert	216,597	73,343	289,940	2.80	5.01	3.15
43	Cultural Expert	123,960	18,037	141,997	1.60	1.23	1.54
44	Activist	97,264	25,189	122,453	1.26	1.72	1.33
45	Health Expert	47,114	7,428	54,542	0.61	0.51	0.59
46	Social Scientist	32,278	19,609	51,887	0.42	1.34	0.56
47	Security Expert	19	8,916	8,935	0.00	0.61	0.10
48	Legal Expert	7	5,025	5,032	0.00	0.34	0.05
<i>Events</i>							

Continued on next page

Table A10 – *Continued from previous page*

#	Category	Count			Proportion (%)		
		EN	FR	ALL	EN	FR	ALL
49	Event Detection	1,421,264	277,560	1,698,824	18.35	18.96	18.44
50	Meeting/Conference	397,039	94,979	492,018	5.13	6.49	5.34
51	Policy Announcement	380,183	42,444	422,627	4.91	2.90	4.59
52	Publication	234,409	72,729	307,138	3.03	4.97	3.33
53	Extreme Weather Event	200,102	38,746	238,848	2.58	2.65	2.59
54	Election	64,663	28,706	93,369	0.83	1.96	1.01
55	Judiciary Decision	37,039	6,285	43,324	0.48	0.43	0.47
56	Cultural Event	964	22,998	23,962	0.01	1.57	0.26
57	Protest	7,741	5,293	13,034	0.10	0.36	0.14
<i>Solutions</i>							
58	Solutions Detection	1,822,433	307,710	2,130,143	23.52	21.02	23.13
59	Mitigation Strategy	1,261,603	201,352	1,462,955	16.29	13.76	15.88
60	Adaptation Strategy	48,963	18,119	67,082	0.63	1.24	0.73
<b>EMOTIONAL TONE</b>							
61	Neutral Emotion	4,879,832	1,057,425	5,937,257	62.99	72.25	64.46
62	Negative Emotion	2,975,222	513,870	3,489,092	38.41	35.11	37.88
63	Positive Emotion	634,048	184,697	818,745	8.18	12.62	8.89
<b>GEOGRAPHIC FOCUS</b>							
64	Canadian Context	3,387,282	573,263	3,960,545	43.72	39.17	43.00
<b>URGENCY/ALARMISM</b>							
65	Urgency/Alarmism	160,220	35,784	196,004	2.07	2.44	2.13
<b>Total Sentences</b>		<b>7,746,839</b>	<b>1,463,581</b>	<b>9,210,420</b>	<b>100.00</b>	<b>100.00</b>	<b>100.00</b>

40

Note: Counts represent the number of sentences in the CCF database annotated with each category. Proportions show the percentage of sentences containing each annotation. Categories are sorted by overall (ALL) prevalence within their hierarchical grouping. EN = English corpus (7.7M sentences), FR = French corpus (1.5M sentences), ALL = Combined corpus (9.2M sentences).

## References

- Do, Salomé, Étienne Ollion, and Rubing Shen (Dec. 2022). “The Augmented Social Scientist: Using Sequential Transfer Learning to Annotate Millions of Texts with Human-Level Accuracy”. In: *Sociological Methods & Research*, p. 00491241221134526. ISSN: 0049-1241. DOI: [10.1177/00491241221134526](https://doi.org/10.1177/00491241221134526). (Visited on 04/15/2024).
- Lemor, Antoine (June 2025). *Antoinelemor/AugmentedSocialScientistFork*. (Visited on 06/19/2025).